

---

# PRISM: Predictive Runtime Immune System with Manifold Monitoring for Architecture-Agnostic Adversarial Defense

---

Rayan Rizon  
Independent Researcher  
rayan.rizon@gmail.com

## Abstract

We present PRISM (Predictive Runtime Immune System with Manifold Monitoring), an architecture-agnostic adversarial defense system that detects adversarial examples at inference time using topological data analysis (TDA) of internal neural network activations. PRISM integrates four complementary components: (1) TAMM (Topological Activation Manifold Monitor), which extracts persistence diagrams from per-layer activations and scores each inference against a calibrated clean-data profile using Wasserstein distance; (2) CADG (Conformal Adversarial Detection Guarantee), which provides distribution-free false-positive rate certificates at three severity tiers via split conformal prediction; (3) SACD (Sequential Adversarial Campaign Detector), a CUSUM-on-score-stream monitor that flags sustained attack campaigns at a mean latency of **2.6 queries** on  $\rho=1.0$  streams while maintaining 0.0 false-alarm activation on clean-only streams; and (4) TAMSH (Topology-Aware MoE Self-Healing), which recovers correct predictions on L3-rejected adversarials through attack-specialist expert sub-networks. The *learned* router (combined  $H_0+H_1$  Wasserstein gating) is the deployable contribution and delivers **+10.6pp** recovery gap over passthrough under PGD — a positive but partial result against a pre-registered +15 pp target — and is isolated from expert capacity by a uniform  $1/K$  control (paired gap 0 pp). A prior-augmented variant that always routes to the PGD specialist (**+26.0pp**) is reported as an *upper-bound analysis* against a per-input oracle ceiling of +36.6 pp, not as a deployable policy.

We evaluate PRISM on three settings (5 seeds,  $n=1000/\text{seed}$ ; for WRN-28-10, additional attack-config repeats are pooled per Table 3): (i) CIFAR-10 with a CIFAR-native ResNet-18 backbone, (ii) CIFAR-10 with a WRN-28-10 backbone, and (iii) CIFAR-100 with a ResNet-18 backbone, under FGSM, PGD, CW-L2, Square, AutoAttack, and adaptive PGD, with per-channel ablation, design-FPR (10%) baseline comparison against LID, Mahalanobis, ODIN, and Energy (matched on CIFAR-10/ResNet-18; three baselines exceed the 10% target on WRN-28-10/CIFAR-100, reported unfiltered), campaign detection, and L3 recovery. On the main CIFAR-10/ResNet-18 setting, PRISM Full achieves 0.918 mean TPR on FGSM/PGD/Square (matched empirical FPR 0.074), 0.938 on CW-L2 (0.075), and 1.000 on AutoAttack (0.075), beating the strongest reproduced baseline (Mahalanobis 0.582 mean TPR at FPR 0.077) by +33.6 pp on the FGSM/PGD/Square pool. The WRN-28-10 backbone lifts the same pool to 0.975 at FPR 0.073; CIFAR-100 holds 0.830 at FPR 0.077. Per-setting tables and the full CIFAR-100 appendix are reported in §4.3 and Appendix N. Unlike detector-coupled defenses, PRISM operates on any pre-trained classifier without retraining or architecture modification, and is, to our knowledge, the first activation-based adversarial detector to provide a split-conformal, distribution-free false-positive-rate certificate.

# 1 Introduction

Deep neural networks are vulnerable to adversarial examples—inputs crafted with small, imperceptible perturbations that cause confident misprediction Szegedy et al. [2014], Goodfellow et al. [2015]. Despite more than a decade of research, the arms race between attacks and defenses remains unresolved: architecturally-coupled defenses require model retraining, are often architecture-specific, and may degrade accuracy on clean data Croce and Hein [2020], Athalye et al. [2018].

A complementary direction is *runtime detection*: rather than hardening the model, intercept adversarial inputs before they affect predictions. Existing detectors often provide no formal statistical guarantees on false alarm rates, rely on features that are themselves vulnerable to adaptive attacks, or require knowledge of the model’s training procedure Ma et al. [2018], Grosse et al. [2017], Lee et al. [2018].

**Key insight.** Clean and adversarial inputs traverse fundamentally different paths through a neural network’s activation manifold. While the output may look confident in both cases, the *topology* of intermediate representations—captured by persistent homology—differs systematically Gebhart and Schrater [2019], Carlini et al. [2019]. PRISM exploits this topological signal at runtime as a computationally tractable anomaly score.

**Contributions.** We make four contributions:

- C1 Topological activation profiler (TAMM).** We compute per-layer persistence features relative to a clean medoid profile and combine them with DCT and softmax-entropy features in a calibrated ensemble. The ensemble-no-TDA ablation measures the marginal value of the topological channel.
- C2 Conformal detection guarantee (CADG).** We apply split conformal prediction to the anomaly score stream, obtaining *distribution-free* FPR certificates at three severity tiers (L1:  $\alpha=0.10$ ; L2:  $\alpha=0.03$ ; L3:  $\alpha=0.005$ ) using only a held-out clean calibration set.
- C3 Temporal campaign detector (SACD).** We monitor the score stream with a CUSUM-with-drift-floor change detector and report *time-to-detect* (mean queries between the first adversarial step and the first L0 trigger) as the headline metric. On sustained  $\rho=1.0$  streams SACD raises the campaign alert in  $2.6 \pm 0.5$  queries while maintaining 0.0 false-alarm activation on clean-only streams—a threat model absent from prior work.
- C4 Routed MoE recovery (TAMSH).** For inputs reaching the L3 severity tier, PRISM routes the input through an attack-specialist expert sub-network rather than issuing a blanket rejection. The learned router (combined  $H_0+H_1$  Wasserstein gating) delivers a +10.6 pp recovery gap over passthrough on PGD-rejected inputs; a uniform  $1/K$  control achieves exactly 0 pp, isolating the contribution to the routing signal rather than expert capacity. A prior-augmented variant that always routes to the PGD specialist reaches +26.0 pp; we report this as an upper-bound analysis against the per-input oracle ceiling of +36.6 pp and discuss the residual gap to a fully learned router in §4.8 as a target for future work.

Empirically, we report locked multi-seed ( $5 \times n=1000$ ) artifact results across three settings — CIFAR-10/ResNet-18, CIFAR-10/WRN-28-10, and CIFAR-100/ResNet-18 — covering detection (FGSM, PGD, Square, CW, AutoAttack, adaptive PGD), four reproduced baselines (LID, Mahalanobis, ODIN, Energy), campaign streams, recovery, and a per-channel ablation that exposes each feature family’s attack-specific failure mode. On the main CIFAR-10/ResNet-18 setting, PRISM Full outperforms the strongest baseline (Mahalanobis) by +33.6 pp mean TPR at matched FPR; WRN-28-10 holds a +16.5 pp gap over Mahalanobis and CIFAR-100 retains +42.6 pp over ODIN (Table 4). The full CIFAR-100 per-component breakdown is in Appendix N.

**Scope and limitations.** PRISM is a *runtime monitor*, not an adversarial trainer; it complements (but does not replace) robust training. Detection latency scales with the number of monitored layers and the TDA subsample size; the submitted results report measured GPU latency from the same artifact pipeline used for detection. Like all detectors, PRISM does not provide a guarantee against adaptive adversaries who are aware of the topological scoring mechanism.

## 2 Related Work

**TDA for adversarial detection.** Gebhart and Schrater [2019] were the first to apply persistent homology to adversarial detection, computing topological features of neuron activations to distinguish clean from adversarial inputs. Relative to that line, PRISM’s novelty lies in using TDA as the backbone of a *full defense pipeline* with formal FPR certificates, campaign-level temporal detection, and expert recovery—not as a standalone detector.

**Conformal prediction for robustness.** Gendler et al. [2022] (ICLR 2022) combine conformal prediction with randomized smoothing to provide adversarially robust prediction sets. Our application of conformal prediction is orthogonal: we use it to calibrate a TDA-based *detector* (not a predictor), providing distribution-free FPR guarantees on the anomaly score threshold via the standard split-conformal coverage theorem Vovk et al. [2005], Angelopoulos and Bates [2023].

**Runtime adversarial detectors.** Feature squeezing Xu et al. [2018] applies input transformations and flags inconsistencies. Mahalanobis detector Lee et al. [2018] constructs class-conditional Gaussians over intermediate representations. LID Ma et al. [2018] uses local intrinsic dimensionality. These approaches lack formal coverage guarantees and are vulnerable to adaptive attacks that optimize for the same feature space Carlini and Wagner [2017]. PRISM combines topological features (harder to optimize against) with conformal calibration (providing formal guarantees).

**Architecture-coupled defenses.** Deep  $k$ -nearest-neighbor (DkNN) style defenses replace deep network layers with topology-aware classifiers in feature space Papernot and McDaniel [2018]; unlike these, PRISM wraps any pre-trained model without modifying its architecture or requiring re-training. Input purification methods Shi et al. [2021], Yoon et al. [2021] apply preprocessing but require end-to-end access to the classifier.

**Temporal adversarial defense.** Session-level adversarial campaign *detection* in neural network inference pipelines is underexplored relative to per-input detection. Classical Bayesian Online Changepoint Detection Adams and MacKay [2007] and CUSUM-style sequential testing have a long history in statistical process control but, to our knowledge, have not been applied to runtime adversarial campaign monitoring in deep classifiers; SACD combines both into a low-latency campaign detector with a strictly positive CUSUM drift floor to bound clean-only false-alarm activation.

## 3 Method

Figure 1 shows the PRISM pipeline. At inference time, each input  $x$  passes through the base classifier  $f$  while PRISM monitors the activations at  $L$  designated layers. The defense proceeds in five stages.

**Notation.** Let  $f : \mathcal{X} \rightarrow \mathbb{R}^C$  be a pre-trained classifier. Denote by  $\phi_\ell(x) \in \mathbb{R}^{d_\ell}$  the activation tensor at layer  $\ell \in \{1, \dots, L\}$ . A *persistence diagram*  $\text{Dgm}(\cdot)$  is a multiset of  $(b_i, d_i)$  pairs encoding topological features (birth/death of connected components in  $H_0$  and loops in  $H_1$ ). The  $p$ -Wasserstein distance between diagrams is

$$W_p(\mathcal{D}_1, \mathcal{D}_2) = \left( \inf_{\gamma} \sum_i \|\gamma(i) - i\|^p \right)^{1/p},$$

where  $\gamma$  ranges over bijections between the points (including diagonal projections for unmatched points) Edelsbrunner and Harer [2010].

### 3.1 TAMM: Topological Activation Manifold Monitor

**Reference profile construction.** Given a held-out profile split  $\mathcal{D}_{\text{profile}} = \{x_i\}_{i=1}^N$  of clean inputs, we extract activations  $\phi_\ell(x_i)$  for each layer  $\ell$  and subsample deterministically to at most  $n_{\text{sub}} = 150$  points to make persistent homology computation tractable. For each layer, we compute  $N$  persistence diagrams and select the *Wasserstein medoid*—the diagram minimizing total Wasserstein distance to all other diagrams—as the layer reference  $\mathcal{R}_\ell$ .




Figure 1: PRISM defense pipeline. Activations from  $L$  layers are converted to persistence diagrams (TMM), scored against a clean reference profile (CADG), monitored for campaigns (SCD), and routed through expert sub-networks if L3 is triggered (TAMSH).

**Anomaly scoring.** At inference time, for input  $x$ , we compute a base topological score

$$s_\ell(x) = W_2(\text{Dgm}(\phi_\ell(x)), \mathcal{R}_\ell), \quad (1)$$

$$S_{\text{TDA}}(x) = \sum_{\ell=1}^L w_\ell s_\ell(x), \quad (2)$$

with live weights  $w_{\text{layer2}} = 0.30$ ,  $w_{\text{layer3}} = 0.30$ , and  $w_{\text{layer4}} = 0.40$ . The publishable detector is not the Wasserstein score alone: PRISM feeds six feature families into a logistic ensemble, then fuses the centered ensemble logit with  $S_{\text{TDA}}(x)$ . Concretely, each input  $x$  produces a feature vector

$$\mathbf{f}(x) = [\mathbf{f}_{\text{TDA}}(x), f_{\text{DCT}}(x), f_{\text{ent}}(x), f_{\text{grad}}(x), \mathbf{f}_{\text{stab}}(x), \mathbf{f}_{\text{logit}}(x)] \in \mathbb{R}^{d_{\text{ens}}},$$

combining six channels:

- **TDA features  $\mathbf{f}_{\text{TDA}}$** : per-layer persistence statistics derived from  $\text{Dgm}(\phi_\ell(x))$  (number of  $H_0/H_1$  features, total persistence, max persistence, persistence entropy).
- **DCT high-frequency energy  $f_{\text{DCT}}(x)$** :  $\sum_{(u,v) \in \mathcal{H}} |X_{u,v}|^2$ , where  $X$  is the 2-D DCT of  $x$  and  $\mathcal{H}$  indexes the upper  $3/4$  band-limited region.
- **Softmax entropy  $f_{\text{ent}}(x)$** :  $-\sum_c p_c(x) \log p_c(x)$  with  $p(x) = \text{softmax}(f(x))$ .
- **Gradient norm  $f_{\text{grad}}(x)$** :  $\|\nabla_x \mathcal{L}_{\text{CE}}(f(x), \hat{y})\|_2$ , computed via `torch.autograd.grad` on the input (not on model parameters), where  $\hat{y} = \arg \max f(x)$ .
- **StabilityV2 features  $\mathbf{f}_{\text{stab}}(x)$** : softmax-confidence drop under  $M=5$  jittered copies  $x + \xi_i$ ,  $\xi_i \sim \mathcal{N}(0, \sigma^2 I)$ , summarised by mean / max / variance of the per-sample confidence deltas. There is no closed-form differentiable surrogate; this is the channel that the present BPDA adaptive attack (§4.4) does not target.
- **Logit-profile features  $\mathbf{f}_{\text{logit}}(x)$** : moments of  $f(x)$  (max-margin, top-2 gap, entropy of  $|f(x)|$ ), capturing classifier-confidence anomalies orthogonal to TDA.

Concretely, the raw feature vector has  $d_{\text{ens}} = 55$  dimensions: 36 TDA persistence statistics (6 statistics  $\times$  2 homology dimensions  $\times$  3 monitored layers), plus  $f_{\text{DCT}}$ ,  $f_{\text{ent}}$ , 8 logit-profile features, 8 StabilityV2 features, and 1 gradient-norm feature. The deployed contract is recorded as `feature_space_version = pixel-stability-v2+logitprofile+sidequad+gradnorm`. The linear heads additionally consume a *side-quadratic* interaction basis: the 190 pairwise products (upper triangle including the diagonal) of the 19 non-TDA side channels (raw indices 36–54), expanding the logistic-regression input to  $55 + 190 = 245$  dimensions. Persistence extraction, the external 55-feature contract, and the split-conformal calibration all operate on the raw 55 features; the 245-dimensional expansion affects only the linear head. The ensemble logit is the logistic regression  $\ell_{\text{ens}}(x) = \mathbf{w}^\top \psi(\mathbf{f}(x)) + b$ , where  $\psi$  appends the side-quadratic terms; it is centred on the calibration

distribution to yield  $\tilde{\ell}_{\text{ens}}(x) = \ell_{\text{ens}}(x) - \ell_{\text{shift}}$ . The final PRISM score is the convex combination  $S(x) = \alpha S_{\text{TDA}}(x) + (1 - \alpha) \tilde{\ell}_{\text{ens}}(x)$  fed to the CADG conformal calibrator, with  $\alpha$  controlling the TDA-vs-ensemble trade-off. The deployed configuration uses  $\alpha = 0.2$  for both CIFAR-10/100 ResNet-18 and  $\alpha = 0.5$  for WRN-28-10 (re-tuned per backbone on the same calibration split). The *ensemble-no-TDA* ablation sets  $\alpha = 0$  (equivalently, zeros  $\mathbf{f}_{\text{TDA}}$  and skips the TDA term) to isolate the marginal value of persistence features. The *channel-mask* ablations (§4.5) zero one of  $\{f_{\text{DCT}}, f_{\text{grad}}, \mathbf{f}_{\text{stab}}, \mathbf{f}_{\text{logit}}\}$  at scoring time while reusing the Full-PRISM threshold, exposing each family’s per-attack contribution.

### 3.2 CADG: Conformal Adversarial Detection Guarantee

PRISM applies split conformal prediction Vovk et al. [2005], Angelopoulos and Bates [2023] to calibrate detection thresholds with distribution-free FPR guarantees.

**Calibration.** Given held-out clean scores  $s_1, \dots, s_n$  (iid from the clean distribution), the conformal threshold at level  $\alpha \in (0, 1)$  is:

$$\hat{q}_\alpha = s_{(\lceil (n+1)(1-\alpha) \rceil)}, \quad (3)$$

where  $s_{(k)}$  denotes the  $k$ -th order statistic.

**Proposition 1** (Coverage guarantee). *For any clean test input  $x_{\text{test}}$  drawn iid from the same distribution as the calibration set,*

$$\Pr[S(x_{\text{test}}) > \hat{q}_\alpha] \leq \alpha + \frac{1}{n+1}.$$

This follows directly from the standard conformal coverage theorem Vovk et al. [2005]. PRISM uses three published FPR tiers with  $\alpha \in \{0.10, 0.03, 0.005\}$ . Thresholds are generated by the artifact pipeline ( $n_{\text{cal}} = 2,000$  clean calibration scores; reproducible via the seed list and conformal quantile in Eq. (3)) and imported into the paper tables; fixed numeric thresholds are not hand-transcribed. All reported TPR/FPR confidence intervals are Wilson 95% binomial intervals computed from pooled per-seed ( $TP, FP, FN, TN$ ) counts; per-seed standard deviations are the sample standard deviation across the five fixed seeds.

**Eval-pool overlap (CI caveat).** Each seed subsamples  $n=1000$  adversarials from the shared 2,000-image evaluation window (`eval_idx`= [8000, 10000]), so the five per-seed pools overlap heavily: a pooled  $n=5000$  count covers only  $\approx 1929$  *unique* source images. The pooled Wilson interval therefore treats overlapping draws as independent and is mildly anti-conservative (too narrow); we report the per-seed standard deviation alongside every pooled interval as the seed-level dispersion check, and flag disjoint per-seed eval windows (or image-clustered CIs) as the clean statistical fix for future runs.

| Tier | $\alpha$ | Threshold source         | Action                         |
|------|----------|--------------------------|--------------------------------|
| L1   | 0.10     | split conformal artifact | Log + normal inference         |
| L2   | 0.03     | split conformal artifact | Input purification + inference |
| L3   | 0.005    | split conformal artifact | Expert routing (MoE) or reject |

Table 1: Conformal detection tiers. Thresholds are calibrated on the active dataset’s clean calibration split and verified on a disjoint validation split.

**Campaign-adaptive thresholds.** When the L0 campaign detector (§3.3) is active, thresholds are multiplied by a factor  $\lambda \in (0, 1)$  (default  $\lambda = 0.8$ ) to increase sensitivity during sustained attacks, accepting a controlled FPR increase.

### 3.3 SACD: Sequential Adversarial Campaign Detection

Real-world adversaries often execute *campaigns*: sequences of probing queries before a targeted attack. SACD monitors the score stream  $S(x_1), S(x_2), \dots$  online via a *centred-score CUSUM* detector whose statistics are computed on top of a Bayesian Online Changepoint Detection (BOCPD) Adams and MacKay [2007] substrate (Normal-Inverse-Gamma conjugate prior, hazard rate 1/30).

BOCPD maintains the posterior run-length distribution  $p(r_t \mid S_{1:t})$  and yields, at each step  $t$ , the short-run probability  $\Pr[r_t \leq k_{\text{alert}}]$  that is available as an alternative trigger:

$$\Pr[r_t \leq k_{\text{alert}}] > p_{\text{alert}}, \quad t > t_{\text{warmup}}, \quad (4)$$

with  $k_{\text{alert}} = 10$ ,  $p_{\text{alert}} = 0.60$ . A grid search over BOCPD-only, CUSUM-only, and hybrid modes (`scripts/calibrate_10_thresholds.py`) selected CUSUM-only as the deployed mode under the constraint *clean-only L0 activation fraction*  $\leq 0.01$ ; the calibrated artifact (`models/10_thresholds.pkl`) is the threshold source for every campaign run reported in this paper.

**Deployed CUSUM trigger.** The deployed detector tracks  $C_t = \max(0, C_{t-1} + (z_t - \delta))$  over centred scores  $z_t = (S(x_t) - \mu_{\text{cal}})/\sigma_{\text{cal}}$  standardised by the calibration-set clean-score mean  $\mu_{\text{cal}} = 0.281$  and standard deviation  $\sigma_{\text{cal}} = 1.639$ . We require a strictly positive drift floor: with  $\delta = 0$ ,  $C_t$  accumulates without the negative reset term and crosses any finite trigger on long enough clean streams, producing up to 59.4% active-fraction on clean-only inputs at seed-42 in our pre-fix runs. The F1 calibration raises  $\delta$  from 0 to **1.0** (one calibration-clean standard deviation) and requires  $k_{\text{consec}}$  consecutive trigger steps before raising the alert, with cusum threshold  $C^* = 5.0$  (CIFAR-10) and warmup  $t_{\text{warmup}} = 20$ . With  $\delta = 1.0$ ,  $k_{\text{consec}} = 3$  on CIFAR-10 (respectively  $k_{\text{consec}} = 5$  on CIFAR-100, calibrated per-dataset), the active fraction on a 500-step clean-only stream is 0.000 on all 5 seeds. Sensitivity to each hyperparameter is reported in Appendix D.

**Headline metric.** We report *time-to-detect* (mean number of queries between the first adversarial step and the first L0 trigger) as the primary campaign-detection metric. A naive “ASR gap (L0 off vs. L0 on)” metric at  $\rho=1.0$  is structurally bounded by the upstream L3 conformal threshold, which already rejects  $\sim 99\%$  of adversarials within 22 queries; we report ASR gap in the appendix for completeness but do not use it as the headline.

### 3.4 TAMSH: Topology-Aware MoE Self-Healing

Inputs that breach the L3 threshold are highly likely adversarial. Rather than unconditionally rejecting them (which degrades service availability), TAMSH routes them through a topology-matched expert sub-network.

**Expert training.** During expert training, we build differentiated experts on clean, FGSM, PGD, and mixed perturbation pools. Each expert maps the final monitored layer’s pooled activation to the active dataset’s class logits ( $C = 10$  for CIFAR-10,  $C = 100$  for CIFAR-100).

**Expert routing.** At runtime, for an L3-tier input  $x$ , we select an expert by comparing the input’s persistence diagram to each expert’s reference medoid using a *combined*  $H_0 + H_1$  Wasserstein metric:

$$k^*(x) = \arg \min_{k=1}^K \left[ w_0 W_2(\text{Dgm}_{H_0}(\phi_L(x)), \mathcal{R}_k^{L, H_0}) + w_1 W_2(\text{Dgm}_{H_1}(\phi_L(x)), \mathcal{R}_k^{L, H_1}) \right], \quad (5)$$

with  $w_0 = w_1 = 1$  in the deployed configuration. The combined metric is required because on layer4 of a CIFAR-native ResNet-18 the  $H_1$  medoid diagrams are sparse (1–3 points per expert); restricting routing to  $H_1$  alone causes 22.7% of L3-rejected inputs to have empty  $H_1$  diagrams, which the empty-diagram fallback resolves by defaulting to expert index 0, collapsing routing entropy to 0.336/2.000 bits. Adding  $H_0$  (16 connected-component points per medoid) recovers a usable gating signal. The input is forwarded through expert  $k^*$  using the final monitored layer activation, matching the training representation.

**Routing variants and the uniform-router control.** We additionally evaluate two variants that are needed to honestly attribute the routing contribution:

- **Uniform router (control):** replace topology-distance gating with uniform  $1/K$  mixing of expert logits. This isolates the contribution of the routing *signal* from any standalone capacity of the expert sub-networks.
- **Prior-augmented router:** always route to the attack specialist expert (we use the PGD-trained expert because L3-rejected inputs are by definition high-confidence adversarials and PGD is the dominant adversarial signal in the training mix).

§4.8 reports recovery accuracy under all three router configurations. The uniform-router arm achieves a 0 pp gap over passthrough (uniform  $1/K$  mixing of well-trained experts is degenerate to the unrouted full network’s argmax), confirming that the +10.6 pp gap of the combined-metric router is attributable to the gating signal, and the further +15.4 pp lift of the prior-augmented router reflects correct expert specialization under a strong prior.

## 4 Experiments

We evaluate PRISM on CIFAR-10 against white-box, black-box, optimization-based, and adaptive attacks, followed by per-channel ablation, matched-FPR baseline comparison, campaign-detection latency, and L3 recovery. To assess cross-architecture and cross-dataset generalization, we additionally report the same detection and baseline protocol on CIFAR-10/WRN-28-10 and CIFAR-100/ResNet-18 (§4.3, Table 3); a full per-table CIFAR-100 set (adaptive PGD, ablation, recovery, campaign, latency) is deferred to Appendix N.

### 4.1 Experimental Setup

**Base classifier.** The main CIFAR-10 results use a CIFAR-native ResNet-18 He et al. [2016] trained from scratch. For cross-architecture validation we additionally train a WRN-28-10 Zagoruyko and Komodakis [2016] on CIFAR-10; for cross-dataset validation we use ResNet-18 on CIFAR-100. PRISM monitors activations at  $L = 3$  residual block outputs weighted 0.30, 0.30, 0.40 (shallow-to-deep) in all settings. For the ResNet-18 backbones (CIFAR-10 and CIFAR-100) these are `layer2`, `layer3`, `layer4`; for the WRN-28-10 backbone (which has three residual stages, not four) the analogous depth slice is `layer1`, `layer2`, `layer3` with the same 0.30/0.30/0.40 weights, preserving the shallow-to-deep emphasis. All inputs remain native  $3 \times 32 \times 32$  images with dataset-specific CIFAR normalization; no ImageNet pretraining or  $224 \times 224$  upsampling is used.

**PRISM calibration.** We calibrate thresholds using  $n_{\text{cal}} = 2,000$  held-out clean images (split conformal:  $n_{\text{cal}}$  for threshold fitting,  $n_{\text{val}} = 1,000$  for coverage verification), with  $\alpha \in \{0.10, 0.03, 0.005\}$ . Thresholds and validation FPRs are generated by the artifact pipeline and reported from the resulting JSON files rather than manually transcribed.

**Attacks.** We evaluate four attack classes:

- **FGSM** Goodfellow et al. [2015]: single-step gradient attack,  $\varepsilon = 8/255$  ( $\ell_\infty$ ).
- **PGD-40** Madry et al. [2018]: 40-step projected gradient descent,  $\varepsilon = 8/255$ , step size  $\alpha = \varepsilon/4$  ( $\ell_\infty$ ).
- **CW-L2** Carlini and Wagner [2017]: optimization-based attack, canonical configuration `max_iter=100, binary_search_steps=9` (matches Table 2 results `cw_n1000_ms5_seed*`). A weakened-CW run (`max_iter=5, binary_search_steps=5`) is reported in Appendix L for detector-fidelity analysis only.
- **Square Attack** Andriushchenko et al. [2020]: black-box score-based attack, 5000 queries,  $\varepsilon = 8/255$  ( $\ell_\infty$ ).
- **AutoAttack** Croce and Hein [2020]: standard ensemble evaluation.
- **Adaptive PGD**: BPDA-style activation-matching attack with  $\lambda$  sweep, 100 steps, 10 restarts, and EOT verification metadata.

All attacks are implemented via Adversarial Robustness Toolbox (ART) v1.20 Nicolae et al. [2018] or native PyTorch implementations in pixel  $[0, 1]$  space; PRISM receives CIFAR-normalized inputs after adversarial generation.

**Baselines.** We reproduce four matched-FPR baselines on the same splits: **LID** Ma et al. [2018], **Mahalanobis** Lee et al. [2018], **ODIN** Liang et al. [2018], and **Energy** Liu et al. [2020].

### 4.2 Attack Detection Results

**Key observations.** The paper reports only numbers produced by the locked artifact pipeline. Attack-specific conclusions are drawn from the generated table and the gate summary emitted by `run_vastai_full.sh` and `run_vastai_cifar100.sh`; no hand-entered historical values are used.

| Dataset  | Attack                        | TPR (mean $\pm$ std) | 95% CI         | FPR          | $n$              |
|----------|-------------------------------|----------------------|----------------|--------------|------------------|
| CIFAR-10 | FGSM                          | 0.882 $\pm$ 0.004    | [0.872, 0.890] | 0.074        | 5000             |
| CIFAR-10 | PGD-40                        | 0.987 $\pm$ 0.003    | [0.984, 0.990] | 0.074        | 5000             |
| CIFAR-10 | Square                        | 0.886 $\pm$ 0.011    | [0.877, 0.894] | 0.074        | 5000             |
| CIFAR-10 | CW-L2                         | 0.938 $\pm$ 0.008    | [0.931, 0.945] | 0.075        | 5000             |
| CIFAR-10 | AutoAttack                    | 1.000 $\pm$ 0.000    | [0.999, 1.000] | 0.075        | 5000             |
| CIFAR-10 | <i>Macro-avg over attacks</i> | <i>0.939</i>         | —              | <i>0.074</i> | <i>5 attacks</i> |

Table 2: Attack detection results generated directly from the Vast.ai result JSONs. TPR is adversarial detection rate; FPR is clean false-positive rate. The final row is the unweighted mean of the five per-attack TPRs (a macro-average across attack families), *not* a pooled-binomial meta-estimate over the 25,000 adversarials; with equal per-attack  $n=5000$  the two point estimates coincide (0.939), so we report the macro-average and claim no pooled confidence interval for it.

### 4.3 Multi-Backbone and Multi-Dataset Validation

To test whether PRISM’s detection generalizes beyond the main CIFAR-10 ResNet-18 setting, we re-ran the full detection and baseline protocol on (i) CIFAR-10 with a WRN-28-10 backbone and (ii) CIFAR-100 with a ResNet-18 backbone, using the same calibration scheme ( $\alpha \in \{0.10, 0.03, 0.005\}$ ) and the same five seeds ( $n = 1000$  adversarials per attack per seed).

| Dataset   | Backbone  | Attack     | TPR   | 95% CI         | FPR   | $n_{\text{adv}}$ |
|-----------|-----------|------------|-------|----------------|-------|------------------|
| CIFAR-10  | ResNet-18 | FGSM       | 0.882 | [0.872, 0.890] | 0.074 | 5000             |
| CIFAR-10  | ResNet-18 | PGD-40     | 0.987 | [0.984, 0.990] | 0.074 | 5000             |
| CIFAR-10  | ResNet-18 | Square     | 0.886 | [0.877, 0.894] | 0.074 | 5000             |
| CIFAR-10  | ResNet-18 | CW-L2      | 0.938 | [0.931, 0.945] | 0.075 | 5000             |
| CIFAR-10  | ResNet-18 | AutoAttack | 1.000 | [0.999, 1.000] | 0.075 | 5000             |
| CIFAR-10  | WRN-28-10 | FGSM       | 0.985 | [0.982, 0.987] | 0.073 | 7000             |
| CIFAR-10  | WRN-28-10 | PGD-40     | 0.995 | [0.993, 0.996] | 0.073 | 7000             |
| CIFAR-10  | WRN-28-10 | Square     | 0.944 | [0.938, 0.949] | 0.076 | 6000             |
| CIFAR-10  | WRN-28-10 | CW-L2      | 0.933 | [0.926, 0.940] | 0.079 | 5000             |
| CIFAR-10  | WRN-28-10 | AutoAttack | 1.000 | [1.000, 1.000] | 0.075 | 10000            |
| CIFAR-100 | ResNet-18 | FGSM       | 0.835 | [0.825, 0.845] | 0.077 | 5000             |
| CIFAR-100 | ResNet-18 | PGD-40     | 0.984 | [0.980, 0.987] | 0.077 | 5000             |
| CIFAR-100 | ResNet-18 | Square     | 0.672 | [0.659, 0.685] | 0.077 | 5000             |
| CIFAR-100 | ResNet-18 | CW-L2      | 0.687 | [0.674, 0.700] | 0.073 | 5000             |
| CIFAR-100 | ResNet-18 | AutoAttack | 0.999 | [0.998, 1.000] | 0.077 | 5000             |

Table 3: Cross-architecture and cross-dataset detection. All three settings clear the L1 design FPR ( $\leq 0.10$ ); per-tier FPR (L2/L3) values are reported in Appendix N and the per-seed JSONs. WRN-28-10 lifts CIFAR-10 TPR by +5–10 pp over ResNet-18 across FGSM/Square; CIFAR-100 is harder for score-based and optimisation-based attacks (Square 0.672, CW 0.687) while white-box PGD-40 and AutoAttack remain at  $\geq 0.98$ . The  $n_{\text{adv}}$  column reports the pooled adversarial count across all 5 seeds and (where applicable) attack-config repeats. CIFAR-10/ResNet-18 and CIFAR-100/ResNet-18 use one config per seed ( $5 \times 1000 = 5000$ ). For WRN-28-10, FGSM/PGD pool the `fgsm_pgd_l3fix_wrn` and `paper_wrn` runs ( $5 + 2$  seed-repeats = 7000), Square pools `prism_square_wrn` and `paper_wrn` ( $4 + 2 = 6000$ ), and AutoAttack pools the `fast_wrn` run plus the `paper_wrn` 5-seed sweep ( $5 + 5 = 10000$ ); CW uses the canonical 5-seed run (5000).

**Cross-setting take-aways.** (i) PRISM transfers to a deeper architecture (WRN-28-10) without re-tuning: the WRN run lifts every CIFAR-10 attack TPR while holding FPR at 0.073, indicating that the discriminative signal is in the ensemble channels, not in a ResNet-18-specific calibration. (ii) Cross-dataset to CIFAR-100 ( $10\times$  more classes), white-box gradient attacks remain trivially detectable (PGD-40 0.984, AutoAttack 0.999), but score-based Square (0.672) and the optimisation-based CW-L2 (0.687) drop below the 0.85 design target. Per-channel ablation (Appendix N, Table 22) shows the gap is concentrated in the StabilityV2 channel, suggesting per-dataset retuning of stability features as a natural next step. (iii) Even where PRISM underperforms its 0.85 target on CIFAR-100,



| Setting              | Detector     | FGSM TPR     | PGD TPR      | Square TPR   | Mean TPR     | Mean FPR     |
|----------------------|--------------|--------------|--------------|--------------|--------------|--------------|
| CIFAR-10, ResNet-18  | LID          | 0.277        | 0.624        | 0.170        | 0.357        | 0.076        |
|                      | Mahalanobis  | 0.393        | 0.877        | 0.476        | 0.582        | 0.077        |
|                      | ODIN         | 0.704        | 0.001        | 0.859        | 0.521        | 0.088        |
|                      | Energy       | 0.595        | 0.001        | 0.974        | 0.523        | 0.105        |
|                      | SID          | 0.296        | 0.519        | 0.286        | 0.367        | 0.082        |
|                      | <b>PRISM</b> | <b>0.882</b> | <b>0.987</b> | <b>0.886</b> | <b>0.918</b> | <b>0.074</b> |
| CIFAR-10, WRN-28-10  | LID          | 0.703        | 0.788        | 0.227        | 0.573        | 0.105        |
|                      | Mahalanobis  | 0.930        | 0.995        | 0.506        | 0.810        | 0.122        |
|                      | ODIN         | 0.750        | 0.000        | 0.867        | 0.539        | 0.104        |
|                      | Energy       | 0.622        | 0.000        | 0.985        | 0.536        | 0.100        |
|                      | SID          | 0.330        | 0.621        | 0.238        | 0.396        | 0.094        |
|                      | <b>PRISM</b> | <b>0.985</b> | <b>0.995</b> | <b>0.944</b> | <b>0.975</b> | <b>0.073</b> |
| CIFAR-100, ResNet-18 | LID          | 0.286        | 0.290        | 0.267        | 0.281        | 0.099        |
|                      | Mahalanobis  | 0.053        | 0.994        | 0.063        | 0.370        | 0.093        |
|                      | ODIN         | 0.767        | 0.000        | 0.445        | 0.404        | 0.116        |
|                      | Energy       | 0.663        | 0.000        | 0.336        | 0.333        | 0.104        |
|                      | SID          | 0.150        | 0.231        | 0.082        | 0.154        | 0.109        |
|                      | <b>PRISM</b> | <b>0.835</b> | <b>0.984</b> | <b>0.672</b> | <b>0.830</b> | <b>0.077</b> |

Table 4: Design-FPR ( $\leq 10\%$ ) baseline comparison across all three settings (5 seeds,  $n = 1000$ ). The comparison is at exactly matched FPR only on CIFAR-10/ResNet-18; on WRN-28-10 and CIFAR-100 three baselines exceed the 10% design target (disclosed below) and are reported unfiltered. PRISM is the only detector that holds  $\geq 0.8$  mean TPR on all three settings while keeping FPR below the design target; every baseline collapses on at least one attack in every setting (ODIN/Energy on PGD; LID on Square; Mahalanobis on FGSM/Square on CIFAR-100). The recent unsupervised detector SID Tian et al. [2021] (2021) is the weakest in every setting (mean TPR 0.15–0.40), confirming that PRISM’s margin is not an artifact of comparing only against pre-2021 baselines. **FPR-target disclosure:** in the WRN-28-10 setting, Mahalanobis (0.122) and LID (0.105) exceed the  $\alpha=0.10$  design FPR; in the CIFAR-100 setting, ODIN (0.116) and SID (0.109) exceed it. Re-calibrating these detectors at a tighter conformal  $\alpha=0.073$  (matched to PRISM’s empirical FPR) would lower their reported TPR further; the present table reports the published-protocol values unfiltered for transparency.




Figure 2: Per-attack TPR at the design FPR target ( $\leq 0.10$ ; not all baselines hit it on WRN/CIFAR-100, see Table 4). PRISM (blue) is the only detector that holds  $\geq 0.85$  on at least three of four attacks per dataset; ODIN/Energy collapse to 0 on PGD; LID and Mahalanobis collapse on Square or FGSM. AutoAttack column has only PRISM because baselines were not run on AutoAttack.

it still beats the strongest baseline on every CIFAR-100 attack (Square: +22.7 pp vs ODIN; FGSM: +6.8 pp vs ODIN). Figure 2 visualises both points; the steeper L3 trade-off on CIFAR-100 in Figure 3 (subfigure 3b) reflects the larger 100-class label space.

#### 4.4 Adaptive Attack Evaluation

Following Carlini et al. [2019], we evaluate PRISM against a white-box adaptive attacker that crafts perturbations explicitly suppressing the detector’s signals. The canonical adaptive attack used to generate Table 5 is the BPDA-style PGD launched with `run_adaptive_pgd.py`




Figure 3: PRISM three-tier operating points (L1, L2, L3) per attack. AutoAttack and PGD-40 saturate near TPR=1 even at the tight L3 FPR ( $\leq 0.005$ ); Square trades TPR for tighter FPR most steeply, consistent with its score-based attack signature being the weakest of the four in the ensemble feature space.

-through-scorer, whose loss is

$$\mathcal{L} = -\mathcal{L}_{\text{CE}}(x', y_{\text{pred}}) + \lambda \sum_{\ell} \frac{1}{D_{\ell}} \|a_{\ell}(x') - a_{\ell}(x)\|_2 + 0.5 \cdot |\text{HF}_{\text{DCT}}(x') - \text{HF}_{\text{DCT}}(x)|,$$

where  $a_{\ell}(x)$  is the per-sample clean activation at monitored layer  $\ell$  and  $\text{HF}_{\text{DCT}}$  is the differentiable high-frequency DCT-energy proxy. The activation-matching term targets TAMM and the DCT-energy term targets the ensemble scorer’s high-frequency channel, so *two* of PRISM’s six raw feature channels are attacked jointly. The remaining channels — StabilityV2 (a jitter-stability statistic with no closed-form differentiable surrogate), gradient-norm, softmax-entropy, and logit-profile (and the 190 side-quadratic interaction terms built from these non-TDA channels, §3.1) — have no dedicated suppression term in the loss, and we explicitly return to this partial coverage in §5 and Appendix N. We sweep  $\lambda \in \{0, 0.5, 1, 2, 5, 10\}$  at 100 PGD steps with 10 random restarts; runs use the deterministic PRISM forward pass, so EOT averaging is verified as a no-op via 20 resampling passes (eot\_verify\_samples=20 recorded in the per-seed JSON metadata).

| Dataset  | $\lambda$ | TPR   | 95% CI         | FPR   |
|----------|-----------|-------|----------------|-------|
| CIFAR-10 | 0.0       | 0.832 | [0.822, 0.842] | 0.075 |
| CIFAR-10 | 0.5       | 0.796 | [0.785, 0.807] | 0.075 |
| CIFAR-10 | 1.0       | 0.753 | [0.741, 0.765] | 0.075 |
| CIFAR-10 | 2.0       | 0.660 | [0.647, 0.673] | 0.075 |
| CIFAR-10 | 5.0       | 0.490 | [0.476, 0.504] | 0.075 |
| CIFAR-10 | 10.0      | 0.429 | [0.416, 0.443] | 0.075 |

Table 5: Adaptive PGD with activation-matching term ( $\lambda$  sweep, 100 steps, 10 restarts, 5 seeds). TPR degrades monotonically as  $\lambda$  grows, bottoming at 0.429 at  $\lambda=10$  while FPR holds at the calibrated 0.075. The  $\lambda=0$  row (0.832) is *not* a non-adaptive PGD reference: the fixed DCT high-frequency suppression term (coefficient 0.5) is still active at  $\lambda=0$ , making this row a “DCT-only adaptive” attack. The non-adaptive PGD-40 reference is 0.987 (Table 2).

**Honest framing.** As  $\lambda$  grows, an attacker who is willing to spend perturbation budget on suppressing the TAMM activation-matching and DCT-energy channels can substantially reduce PRISM’s CIFAR-10 detection rate (TPR drops from 0.832 at  $\lambda=0$  to 0.429 at  $\lambda=10$ ); FPR is unaffected because the conformal threshold is fixed on clean calibration scores and does not depend on the adversary’s loss. On CIFAR-100 the same attack only moves TPR from 0.950 to 0.925; we attribute this to the StabilityV2 channel, which the present attack does not target and which the ablation (Table 22)

identifies as CIFAR-100’s dominant discriminative channel (Appendix N). A genuinely ensemble-complete adaptive attack — adding differentiable surrogates for the StabilityV2, gradient-norm, and softmax-entropy losses — is the most important open evaluation gap (§5); the present sweep bounds the achievable damage of the two-channel attacker. We do not claim adaptive-attack robustness against an ensemble-complete attacker; instead, we report degradation curves so deployments can choose operating points and so future work can benchmark against a documented baseline.

#### 4.5 Ablation Study

We ablate seven configurations: (1) Full PRISM; (2–5) four channel-mask arms that zero one ensemble feature family at scoring time — No-LogitProfile, No-StabilityV2, No-GradNorm, No-DCT — each isolating a distinct channel’s contribution; (6) Ensemble-no-TDA, which removes persistence features while keeping softmax-entropy, gradient-norm, DCT, and stability channels; and (7) TDA-only with its own base-Wasserstein conformal calibrator. The channel-mask design replaces an earlier "No MoE" arm that was statistically indistinguishable from Full PRISM ( $t=0$ ,  $p=1$  across all attacks) because the MoE router is a downstream gating module that does not alter ensemble scores; the new arms instead surface each feature family’s attack-specific failure mode (Table 6). SACD/L0 is evaluated separately in campaign streams because it is a temporal mechanism, not an IID detector.

| Dataset  | Configuration   | FGSM TPR                            | PGD TPR                             | Square TPR        | Mean TPR     | Mean FPR     |
|----------|-----------------|-------------------------------------|-------------------------------------|-------------------|--------------|--------------|
| CIFAR-10 | Full PRISM      | $0.882 \pm 0.004$                   | $0.987 \pm 0.003$                   | $0.886 \pm 0.011$ | <b>0.918</b> | 0.074        |
| CIFAR-10 | No-LogitProfile | $0.988 \pm 0.002$                   | $0.980 \pm 0.003$                   | $0.920 \pm 0.003$ | 0.963        | <b>0.938</b> |
| CIFAR-10 | No-StabilityV2  | $0.936 \pm 0.004$                   | <b>0.396 <math>\pm</math> 0.014</b> | $0.854 \pm 0.010$ | 0.729        | 0.107        |
| CIFAR-10 | No-GradNorm     | $0.946 \pm 0.008$                   | $0.985 \pm 0.003$                   | $0.993 \pm 0.003$ | 0.975        | 0.218        |
| CIFAR-10 | No-DCT          | <b>0.751 <math>\pm</math> 0.006</b> | $0.995 \pm 0.002$                   | $0.960 \pm 0.003$ | 0.902        | 0.098        |
| CIFAR-10 | Ensemble-no-TDA | <b>0.506 <math>\pm</math> 0.008</b> | $0.969 \pm 0.003$                   | $0.886 \pm 0.007$ | 0.787        | 0.077        |
| CIFAR-10 | TDA only        | $0.009 \pm 0.002$                   | $0.066 \pm 0.007$                   | $0.019 \pm 0.003$ | 0.031        | 0.082        |

Table 6: IID ablation generated from the post-fix multi-seed ablation JSON (5 seeds,  $n = 1000$  per seed). C1 is evaluated as the marginal gap between Full PRISM and Ensemble-no-TDA (mean TPR  $0.918 - 0.787 = +13.1$  pp). Each channel-mask arm exposes a distinct failure mode: No-LogitProfile FPR explodes to 0.938 (the logit channel is the FPR-controller); No-StabilityV2 PGD collapses to 0.396 (stability is the PGD signal); No-GradNorm FPR triples to 0.218 (gradient norm is a counter-evidence channel); No-DCT FGSM drops to 0.751 (DCT is the FGSM signal). All channel arms reach  $p < 0.001$  on at least one attack. **Reading guide (matching Table 22).** Rows whose mean FPR exceeds the 0.10 design target — No-LogitProfile (0.938), No-GradNorm (0.218), and No-StabilityV2 (0.107) — are threshold collapses, not improved detection: their higher mean TPR is bought by flagging far more clean inputs and is non-deployable. Only Full PRISM (0.074), No-DCT (0.098), and Ensemble-no-TDA (0.077) sit at the design FPR, so only their TPRs are comparable to Full PRISM.

**Baseline comparison.** Table 7 reports four baseline detectors at the same  $\alpha=0.10$  design target, aggregated across the same 5 seeds. PRISM Full has both the lowest empirical FPR (0.074) and the highest mean TPR (0.918), outperforming the strongest baseline (Mahalanobis, mean TPR 0.582 at FPR 0.077) by +33.6 pp mean TPR. PRISM is the only detector that does not collapse on at least one attack family (ODIN/Energy: PGD TPR  $\approx 0.001$ ; LID: Square TPR 0.170). Energy’s empirical FPR (0.105) modestly exceeds its  $\alpha=0.10$  design target on this test split, which we report unfiltered; recomputing at a tighter conformal  $\alpha=0.075$  (matched to PRISM’s empirical FPR) would lower Energy’s mean TPR further. We additionally benchmark two recent (2021) detectors—SID Tian et al. [2021] and the supervised SpectralDefense Harder et al. [2021]—in Appendix Q. Both underperform PRISM: SID across all attack families, and SpectralDefense everywhere except the gradient- $L_\infty$  attacks it is trained on, collapsing on black-box (Square) and  $L_2$  (CW) attacks even in-distribution (Table 29) and failing to transfer across attack families (Table 30).

**Temporal and recovery contributions.** SACD/L0 is reported from campaign-stream experiments (§4.7) and TAMSH from L3 recovery experiments (§4.8). Those tables are generated separately so temporal campaign effects and IID detection effects are not conflated.

| Detector                  | FGSM TPR                            | PGD TPR                             | Square TPR                          | Mean TPR      | Mean FPR     |
|---------------------------|-------------------------------------|-------------------------------------|-------------------------------------|---------------|--------------|
| LID                       | $0.277 \pm 0.009$                   | $0.624 \pm 0.016$                   | $0.170 \pm 0.008$                   | 0.357         | 0.076        |
| Mahalanobis               | $0.393 \pm 0.010$                   | $0.877 \pm 0.005$                   | $0.476 \pm 0.015$                   | 0.582         | 0.077        |
| ODIN                      | $0.704 \pm 0.009$                   | $0.001 \pm 0.000$                   | $0.859 \pm 0.005$                   | 0.521         | 0.088        |
| Energy                    | $0.595 \pm 0.004$                   | $0.001 \pm 0.000$                   | $0.974 \pm 0.003$                   | 0.523         | 0.105        |
| <b>PRISM (Full)</b>       | <b><math>0.882 \pm 0.004</math></b> | <b><math>0.987 \pm 0.003</math></b> | <b><math>0.886 \pm 0.011</math></b> | <b>0.918</b>  | <b>0.074</b> |
| $\Delta$ vs best baseline | +0.178                              | +0.110                              | −0.088                              | <b>+0.336</b> | −0.002       |

Table 7: Matched-FPR baseline comparison (5 seeds,  $n = 1000$  per seed). PRISM Full beats the best baseline by +33.6 pp mean TPR while holding the lowest FPR.

#### 4.6 Score Distribution Analysis

**Reporting.** Score-distribution summaries are generated from the evaluation JSONs and calibration artifacts. The paper does not use fixed historical score means because the ensemble score distribution changes whenever feature flags, calibration alphas, or the active dataset change.




Figure 4: Score distributions for clean and adversarial inputs (CIFAR-10, ResNet-18, single-attack reference). Dashed lines: conformal thresholds L1, L2, L3. Clean and PGD-40 distributions are clearly separated; FGSM at low  $\varepsilon$  overlaps the clean tail.

#### 4.7 Campaign Detection (SACD): Latency, Not ASR Gap

We evaluate L0 campaign detection on clean-only, sustained ( $\rho \in \{0.5, 0.8, 1.0\}$ ), low-rate ( $\rho=0.05$ ), and burst stream scenarios generated by `run_campaign_eval.py`.

**Headline metric: time-to-detect.** On sustained  $\rho=1.0$  streams, the L0 monitor flags the campaign within a mean of  $2.6 \pm 0.5$  queries after the first adversarial step (across 5 seeds; Table 8). On  $\rho=0.5$  streams, detection takes  $3.4 \pm 0.9$  queries; on 10-step burst streams,  $12.0 \pm 0.0$  queries; on low-rate  $\rho=0.05$  streams,  $17.4 \pm 8.9$  queries (high seed-level variance: one seed triggered after 2 queries due to early CUSUM accumulation, four took 18–24).

**Why time-to-detect, not ASR gap.** A naive headline metric — adversarial success rate (ASR) reduction with L0 on vs. off at  $\rho=1.0$  — is structurally bounded by the upstream L3 conformal threshold, which already rejects approximately 99% of  $\rho=1.0$  adversarials within 22 queries, leaving no room for L0 to contribute an additional ASR-reduction gap. We therefore report ASR gap in the appendix for completeness, but treat detection latency (time-to-first-trigger after campaign onset) as the operationally meaningful metric: a defender that learns of an attack within 3 queries can deploy




Figure 5: Per-attack score envelopes (p5–p95 whiskers, IQR boxes, medians) of the fused ensemble score  $S(x)$  for clean (blue) vs adversarial (red) inputs across all four attacks per dataset.  $S(x)$  is signed: clean inputs centre near zero (CIFAR-10 clean median  $\approx 0.3$ ) and attacks push it positive. AutoAttack pushes scores furthest from the clean median (CIFAR-10 adversarial median  $\approx 17$ ); Square has the tightest separation, matching the lower Square TPR in Table 3. CIFAR-100 (right) shows the same ordering with a slightly larger clean spread (clean median  $\approx 0.04$ , p5–p95 width 5.9 vs 4.1 on CIFAR-10), reflecting the 100-class softmax temperature. This is the same  $S(x)$  scale as the per-attack histograms in Appendix J (clean mean 0.27).

| Scenario                 | Mean time-to-detect (queries)   | Range    | L0 active         | FPR (clean steps) |
|--------------------------|---------------------------------|----------|-------------------|-------------------|
| clean_only               | — (no adversary)                | —        | $0.000 \pm 0.000$ | $0.084 \pm 0.028$ |
| sustained_ $\rho=0.50$   | $3.4 \pm 0.9$                   | [2, 4]   | $0.893 \pm 0.030$ | $0.107 \pm 0.053$ |
| sustained_ $\rho=0.80$   | <b><math>2.6 \pm 0.5</math></b> | [2, 3]   | $0.897 \pm 0.014$ | $0.092 \pm 0.045$ |
| sustained_ $\rho=1.00$   | <b><math>2.6 \pm 0.5</math></b> | [2, 3]   | $0.897 \pm 0.014$ | $0.090 \pm 0.055$ |
| burst (10-step pulse)    | $12.0 \pm 0.0$                  | [12, 12] | $0.863 \pm 0.018$ | $0.118 \pm 0.052$ |
| low_rate ( $\rho=0.05$ ) | $17.4 \pm 8.9$                  | [2, 24]  | $0.951 \pm 0.020$ | $0.143 \pm 0.027$ |

Table 8: Campaign detection across stream scenarios (5 seeds, post-fix). The headline metric is mean time-to-detect; clean-only L0 activation fraction is 0.000 on all 5 seeds after the F1 CUSUM drift-floor fix.

mitigations (rate-limiting, account quarantine, model rollback) before the attacker has issued more than a handful of queries against the deployed system.

**Clean-only false alarm rate.** After the F1 CUSUM drift-floor fix, L0 activation fraction on the 500-step calibration clean streams is 0.000 on all 5 seeds (pre-fix: up to 0.594 on 2 of 5 seeds, a defect that would not have survived peer review pre-fix). For 50-query clean windows (warmup= 20, 30 at-risk steps), empirical false alarm probability is  $< 5\%$ . A stress test on longer clean streams (bootstrap-resampled from  $n=1,000$  real post-fusion ensemble scores generated end-to-end on the VAL split) confirms that the deployed  $C^* = 5.0$  holds the 0.000 active-fraction through 1,000 steps and starts admitting its first spurious trigger between step  $\approx 1,800$  and step  $\approx 4,500$  on real long streams (Appendix E, Table 12). The alternative feasible cells  $C^* \in \{8.0, 10.0, 12.0\}$  from the calibration grid hold 0.000 across all stream lengths up to 10,000 steps. Longer-window deployments should either pick one of these conservative cells or periodically reset L0 / auto-deactivate on return to the clean score baseline.

**Two distinct false-alarm rates.** The *FPR (clean steps)* column in Table 8 (0.084–0.143) is the per-query L1/L2/L3 conformal alert rate of the underlying IID detector — the same FPR reported in Table 2 — not the L0 campaign-trigger rate. A deployment that quarantines accounts should gate on L0 firing (clean-only rate 0.000 after the F1 fix), not on every L1 step; L1 alerts are appropriate for softer signals such as elevated logging or rate-limiting, where a  $\sim 10\%$  per-query rate is acceptable because the action is reversible and per-query.

#### 4.8 L3 Recovery (TAMSH)

We measure recovery accuracy of five routing strategies applied to the subset of PGD inputs that trigger the L3 (reject) tier: the most aggressive adversarials, where simply rejecting the input is the conservative default. We pre-registered a +15 pp recovery gap (best learned router vs. passthrough) as the contribution target for this component. As reported below, the *learned* router (combined  $H_0+H_1$  Wasserstein gating) falls 4.4 pp short of that target at +10.6 pp; we therefore report this component as a positive-signal contribution rather than a target-clearing result, and document the residual gap to the prior-augmented variant and the oracle ceiling as the path for future learned routing.

| Strategy                                                                                                 | Recovery Acc. | 95% CI                | Gap vs passthrough (pp) | Avail. | Router signal               |
|----------------------------------------------------------------------------------------------------------|---------------|-----------------------|-------------------------|--------|-----------------------------|
| reject                                                                                                   | 0.000         | [0.000, 0.001]        | —                       | 0.000  | —                           |
| passthrough                                                                                              | 0.033         | [0.027, 0.038]        | —                       | 1.000  | none                        |
| tamsh_uniform                                                                                            | 0.033         | [0.027, 0.038]        | +0.00 ± 0.00            | 1.000  | uniform 1/ $K$ (control)    |
| <b>tamsh (combined <math>H_0+H_1</math>, learned)</b>                                                    | <b>0.139</b>  | <b>[0.128, 0.149]</b> | <b>+10.62 ± 1.45</b>    | 1.000  | Wasserstein gating          |
| <i>Upper-bound analysis (not a deployable policy):</i>                                                   |               |                       |                         |        |                             |
| tamsh_force (PGD specialist)                                                                             | 0.293         | [0.280, 0.306]        | +26.00 ± 1.84           | 1.000  | oracle prior (always-route) |
| <i>Oracle ceiling (per-input best-of-<math>K</math> on the same L3 subset): 0.399 recovery accuracy.</i> |               |                       |                         |        |                             |

Table 9: L3-recovery on PGD adversarials (5 seeds,  $n_{L3} \approx 792$  per seed; mean trigger rate 0.792 within sanity band [0.10, 0.82]). The learned-router row (tamsh, combined  $H_0+H_1$ ) is the primary contribution and delivers +10.6 pp over passthrough, falling 4.4 pp short of the pre-registered +15 pp target. The tamsh\_force row is a prior-augmented upper-bound analysis (always-route to PGD specialist, not a learned policy); the oracle ceiling is per-input best-of- $K$  on the same L3 subset.

**Router contribution isolation.** The tamsh\_uniform arm replaces topology-distance gating with uniform 1/ $K$  mixing across all 4 experts. Its recovery accuracy equals that of passthrough (0.033 each, paired gap 0.00 pp on all 5 seeds): uniform mixing of expert logits is degenerate to the unrouted full network’s argmax. The +10.6 pp tamsh gap (combined  $H_0+H_1$  Wasserstein gating) over uniform is therefore attributable entirely to the routing signal, not to standalone expert-network capacity — a positive but partial result against the pre-registered +15 pp target. The further +15.4 pp lift achieved by tamsh\_force (always-route to the PGD-trained expert) reflects correct routing under a strong prior: L3- rejected inputs are by definition high-confidence adversarials, and PGD is the dominant adversarial signal in our training mix. We treat tamsh\_force as an upper-bound analysis rather than a deployable policy, since it bakes in an oracle assumption about attack identity that a real router would have to learn online. The gap between the learned router (0.139 recovery accuracy) and the prior-augmented variant (0.293) bounds the headroom achievable by a stronger gating signal below the oracle ceiling (0.399).

#### 4.9 Calibration Sensitivity

Figure 6 sweeps  $n_{\text{cal}} \in \{500, 1000, 2000, 3000\}$  and confirms threshold stability: the L3 threshold varies by  $\pm 0.03$  (range 0.06) across calibration set sizes (25 bootstrap subsamples per size), and the achieved FPR on the held-out validation split tracks the target  $\alpha$  within the Wilson-95% interval, satisfying Proposition 1.

## 5 Conclusion

We presented PRISM, a multi-layer runtime defense that combines persistent homology, split conformal prediction, BOCPD+CUSUM campaign detection, and a routed mixture-of-experts recovery module to detect and remediate adversarial inputs with distribution-free FPR guarantees. On CIFAR-10 (5 seeds,  $n=1000$ ), PRISM Full achieves mean detection TPR 0.918 at matched FPR 0.074 across FGSM, PGD-40, and Square, 0.938 on canonical CW-L2, and 1.000 on AutoAttack — +33.6 pp mean TPR over the strongest baseline detector (Mahalanobis). On campaign streams SACD flags sustained  $\rho=1.0$  attacks within  $2.6 \pm 0.5$  queries with zero false-alarm activation on clean-only streams. On L3-rejected PGD adversarials, the deployable TAMSH *learned* router (combined  $H_0+H_1$  Wasserstein gating) recovers correct predictions with +10.6 pp over passthrough, falling




Figure 6: Conformal threshold stability vs. calibration set size  $n_{\text{cal}}$ . *Left*: L1/L2/L3 threshold values (mean  $\pm$  95% bootstrap envelope; L3 half-range  $\pm 0.03$  across the full sweep). *Right*: Achieved FPR on the held-out validation set tracks the target  $\alpha$  for all three tiers. Source: `clean_scores.npy` (2,000 calibration scores) with 25 bootstrap subsamples per size.

4.4 pp short of the pre-registered +15 pp target; a uniform  $1/K$  router control achieves 0 pp, isolating the routing-signal contribution *on PGD*. This uniform-is-degenerate result is PGD-specific: on CW and Square the uniform  $1/K$  mixture matches or beats the topology gate (Appendix I, e.g. Square +47.2 vs +43.5 pp, CW +47.0 vs +39.0 pp), so uniform mixing of the deployed experts — not the topology gate — is the deployable cross-attack recovery strategy beyond PGD. A prior-augmented upper-bound variant that always routes to the PGD specialist reaches +26 pp and is reported as an analysis ceiling against the per-input oracle (+36.6 pp), not as a deployable policy. All numbers are generated from the locked CIFAR-10 Vast.ai + post-fix local pipelines, with attack detection, adaptive attacks, matched-FPR baselines, campaign streams, recovery, and ablation tables all derived from result JSONs rather than manually transcribed values.

**Limitations.** The primary limitation is computational: persistent homology over high-dimensional activation tensors requires subsampling (150 points per layer in the current configuration), which may weaken detection for attacks that specifically preserve local topology while altering global structure. Single-sample end-to-end detection latency on an RTX 4090 (Table 26,  $n = 200$ ) is 39.78 ms mean / 41.26 ms p95 on CIFAR-10 and 50.35 ms mean / 55.90 ms p95 on CIFAR-100, both well below the 100 ms real-time target; CPU-only deployment is significantly slower because persistent-homology computation on high-dimensional activation tensors does not vectorise. Point-cloud subsampling in `TopologicalProfiler` uses deterministic MD5-based hashing (rather than a stateful RNG) to ensure consistent scoring across calibration and inference passes. The conformal calibration assumes the clean test distribution matches the calibration distribution—distribution shift (e.g., domain adaptation attacks) may inflate FPR. The adaptive-attack evaluation in §4.4 targets *two* of PRISM’s ensemble channels — per-layer activation matching (TAMM) and DCT high-frequency energy — via the `-through-scorer` run, and is the canonical adaptive result in this paper. It does *not* attack StabilityV2, gradient-norm, softmax-entropy, or logit-profile (nor the side-quadratic interaction terms built from these channels): StabilityV2 has no closed-form differentiable surrogate (it is a jitter-stability statistic), and gradient-norm, softmax-entropy, and logit-profile were not added to the BPDA loss in this campaign. This partial coverage is the most likely reason CIFAR-100 holds 0.925 TPR at  $\lambda=10$  while CIFAR-10 collapses to 0.429: CIFAR-100’s discriminative weight is concentrated in the StabilityV2 channel (Table 22), which the present attack does not touch. An ensemble-complete adaptive attack (StabilityV2 + gradient-norm + softmax-entropy surrogates added to the `-through-scorer` loss) is the highest-priority open evaluation; we do not claim robustness against such an attacker, and the  $\lambda$ -sweep results bound only what the two-channel attacker can achieve.

**Future work.** Highest-priority directions, in order: (1) an ensemble-complete adaptive attack that augments the `-through-scorer` loss with differentiable surrogates for the StabilityV2 (jitter-stability), gradient-norm, softmax-entropy, and logit-profile channels, closing the two-channel gap identified above; (2) extending the monitor to transformer architectures via attention-head and MLP-

block activation patterns; (3) certified robustness integration by combining PRISM with randomized smoothing Cohen et al. [2019] for joint detection-plus-certified prediction guarantees; (4) online recalibration of the conformal thresholds to handle clean-side distribution shift; (5) a stronger learned router for TAMSH (closing the  $+10.6 \rightarrow +26$  pp gap toward the prior-augmented upper bound and the  $+36.6$  pp oracle ceiling) via gating networks trained directly on the L3-rejected pool rather than via Wasserstein medoid distance alone.

**Cross-architecture and cross-dataset generalisation.** Section 4.3 reports the same protocol on CIFAR-10/WRN-28-10 (mean pool TPR 0.975, FPR 0.073) and CIFAR-100/ResNet-18 (mean pool TPR 0.830, FPR 0.077), both run on the same 5 seeds and same calibration scheme. The full per-component CIFAR-100 breakdown (adaptive PGD, channel-mask ablation with FPR disclosure, campaign, recovery, latency) is in Appendix N. The TAMSH C4 recovery contribution does not transfer to the 100-class setting (within 1.4 pp of passthrough) and is reported as appendix-only on CIFAR-100 per the pre-registered demotion clause.

**Broader impact.** PRISM is a defensive tool designed to improve the reliability of ML systems under adversarial conditions. We do not foresee misuse potential beyond standard ML security research.

## References

- Ryan Prescott Adams and David JC MacKay. Bayesian online changepoint detection. *arXiv preprint arXiv:0710.3742*, 2007.
- Maksym Andriushchenko, Francesco Croce, Nicolas Flammarion, and Matthias Hein. Square attack: a query-efficient black-box adversarial attack via random search. In *European Conference on Computer Vision (ECCV)*, pages 484–501, 2020.
- Anastasios N Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv preprint arXiv:2107.07511*, 2023.
- Anish Athalye, Nicholas Carlini, and David Wagner. Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples. In *International Conference on Machine Learning (ICML)*, pages 274–283, 2018.
- Nicholas Carlini and David Wagner. Towards evaluating the robustness of neural networks. In *IEEE Symposium on Security and Privacy (S&P)*, pages 39–57, 2017.
- Nicholas Carlini, Anish Athalye, Nicolas Papernot, Wieland Brendel, Jonas Rauber, Dimitris Tsipras, Ian Goodfellow, Aleksander Madry, and Alexey Kurakin. On evaluating adversarial robustness. In *arXiv preprint arXiv:1902.06705*, 2019.
- Jeremy M Cohen, Elan Rosenfeld, and J Zico Kolter. Certified adversarial robustness via randomized smoothing. In *International Conference on Machine Learning (ICML)*, pages 1310–1320, 2019.
- Francesco Croce and Matthias Hein. Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks. In *International Conference on Machine Learning (ICML)*, pages 2206–2216, 2020.
- Herbert Edelsbrunner and John Harer. *Computational Topology: An Introduction*. American Mathematical Society, 2010.
- Thomas Gebhart and Paul Schrater. Adversary detection in neural networks via persistent homology. In *arXiv preprint arXiv:1910.13467*, 2019.
- Asaf Gendler, Tsui-Wei Weng, Luca Daniel, and Yaniv Romano. Adversarially robust conformal prediction. In *International Conference on Learning Representations (ICLR)*, 2022.
- Ian J Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples. In *International Conference on Learning Representations (ICLR)*, 2015.
- Kathrin Grosse, Praveen Manoharan, Nicolas Papernot, Michael Backes, and Patrick McDaniel. On the (statistical) detection of adversarial examples. In *arXiv preprint arXiv:1702.06280*, 2017.



- GUDHI Project. GUDHI user and reference manual. *GUDHI Editorial Board*, 2015. URL <https://gudhi.inria.fr/doc/latest/>.
- Paula Harder, Franz-Josef Pfreundt, Margarete Keuper, and Janis Keuper. SpectralDefense: Detecting adversarial attacks on CNNs in the Fourier domain. In *International Joint Conference on Neural Networks (IJCNN)*, pages 1–8, 2021.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- Kimin Lee, Kibok Lee, Honglak Lee, and Jinwoo Shin. A simple unified framework for detecting out-of-distribution samples and adversarial attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 31, 2018.
- Shiyu Liang, Yixuan Li, and R. Srikant. Enhancing the reliability of out-of-distribution image detection in neural networks. In *International Conference on Learning Representations (ICLR)*, 2018.
- Weitang Liu, Xiaoyun Wang, John Owens, and Yixuan Li. Energy-based out-of-distribution detection. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Xingjun Ma, Bo Li, Yisen Wang, Sarah M Erfani, Sudanthi Wijewickrema, Grant Schoenebeck, Dawn Song, Michael E Houle, and James Bailey. Characterizing adversarial subspaces using local intrinsic dimensionality. In *International Conference on Learning Representations (ICLR)*, 2018.
- Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. Towards deep learning models resistant to adversarial attacks. In *International Conference on Learning Representations (ICLR)*, 2018.
- Maria-Irina Nicolae, Mathieu Sinn, Minh-Ngoc Tran, Beat Buesser, Amrith Rawat, Martin Wistuba, Valentina Zantedeschi, Nathalie Baracaldo, Bryant Chen, Heiko Ludwig, et al. Adversarial robustness toolbox v1.0.0. *arXiv preprint arXiv:1807.01069*, 2018.
- Nicolas Papernot and Patrick McDaniel. Deep  $k$ -nearest neighbors: Towards confident, interpretable and robust deep learning. In *arXiv preprint arXiv:1803.04765*, 2018.
- Erich Schubert and Peter J Rousseeuw. Fast and eager  $k$ -medoids clustering: runtime improvement of the PAM, CLARA, and CLARANS algorithms. *Information Systems*, 101:101804, 2021.
- Changhao Shi, Chen Holtz, and Gal Mishne. Online adversarial purification based on self-supervised learning. In *International Conference on Learning Representations (ICLR)*, 2021.
- Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. Intriguing properties of neural networks. In *International Conference on Learning Representations (ICLR)*, 2014.
- Jinyu Tian, Jiantao Zhou, Yuanman Li, and Jia Duan. Detecting adversarial examples from sensitivity inconsistency of spatial-transform domain. In *AAAI Conference on Artificial Intelligence (AAAI)*, pages 9877–9885, 2021.
- Vladimir Vovk, Alex Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- Weilin Xu, David Evans, and Yanjun Qi. Feature squeezing: Detecting adversarial examples in deep neural networks. In *Network and Distributed System Security Symposium (NDSS)*, 2018.
- Jongmin Yoon, Sung Ju Hwang, and Juho Lee. Adversarial purification with score-based generative models. In *International Conference on Machine Learning (ICML)*, 2021.
- Sergey Zagoruyko and Nikos Komodakis. Wide residual networks. In *British Machine Vision Conference (BMVC)*, 2016.

## A Proof of Conformal Coverage (Proposition 1)

We reproduce the standard split conformal argument for completeness.

*Proof.* Let  $S_1, \dots, S_n, S_{\text{test}}$  be iid from distribution  $P$  over  $\mathbb{R}$ . Define  $\hat{q}_\alpha = S_{(\lceil (n+1)(1-\alpha) \rceil)}$ , the  $\lceil (n+1)(1-\alpha) \rceil$ -th order statistic of  $S_1, \dots, S_n$ .

By exchangeability of  $(S_1, \dots, S_n, S_{\text{test}})$ , the rank of  $S_{\text{test}}$  among  $S_1, \dots, S_n, S_{\text{test}}$  is uniformly distributed on  $\{1, \dots, n+1\}$ . Therefore:

$$\Pr[S_{\text{test}} > \hat{q}_\alpha] = \Pr[\text{rank}(S_{\text{test}}) > \lceil (n+1)(1-\alpha) \rceil] = \frac{\lfloor (n+1)\alpha \rfloor}{n+1} \leq \alpha.$$

Adding the  $\frac{1}{n+1}$  slack for the finite-sample correction (since the empirical quantile can shift by at most one order statistic) gives the stated bound.  $\square$

In our calibration, the active dataset uses  $n = 2,000$  clean calibration examples, giving finite-sample slack below  $5 \times 10^{-4}$ .

## B TDA Computation Details

**Persistence homology pipeline.** For each layer  $\ell$ , the activation tensor  $\phi_\ell(x) \in \mathbb{R}^{d_\ell}$  is downsampled to a point cloud  $\mathcal{P} \subset \mathbb{R}^{d_\ell}$  with  $|\mathcal{P}| \leq 150$  via deterministic hash-based subsampling. We then:

1. Construct the Vietoris-Rips filtration on  $\mathcal{P}$  with parameter  $r \in [0, r_{\max}]$ ,  $r_{\max} = 50$ .
2. Compute  $H_0$  (connected components) and  $H_1$  (1-cycles) persistence.
3. Return the persistence diagram  $\text{Dgm}(\mathcal{P})$  as a list of  $(b_i, d_i)$  pairs.

We use gudhi v3.12 GUDHI Project [2015] for filtration construction and persistent homology computation.

**Wasserstein medoid computation.** Given  $N$  diagrams  $\{\mathcal{D}_i\}_{i=1}^N$ , the medoid is:

$$i^* = \arg \min_i \sum_{j=1}^N W_2(\mathcal{D}_i, \mathcal{D}_j).$$

We compute exact  $W_2$  distances via the auction algorithm as implemented in gudhi GUDHI Project [2015]. For  $N = 5,000$  profile diagrams with  $\sim 20$  points each, full pairwise computation takes  $\sim 8$  minutes on a single CPU core.

**Expert clustering.**  $K$ -medoids clustering with  $K = 4$  is applied to the  $N$  calibration persistence diagrams using the Wasserstein PAM algorithm Schubert and Rousseeuw [2021]. Each cluster medoid serves as the routing signature  $\mathcal{R}_k^L$  in Eq. (5). Expert sub-networks are 3-layer MLPs (Linear  $\rightarrow$  ReLU  $\rightarrow$  LayerNorm  $\rightarrow$  Linear  $\rightarrow$  ReLU  $\rightarrow$  LayerNorm  $\rightarrow$  Linear) with hidden dimension 256, trained on per-attack-cluster activation pools. We use LayerNorm rather than BatchNorm because the per-cluster pool size at inference can be as small as 1 (a single L3-rejected input).

## C Campaign Detection Latency

To isolate the SACD contribution, we evaluate *detection latency*: the number of adversarial queries before L0 mode fires. We construct a synthetic query stream: 50 clean images, then consecutive PGD/FGSM adversarial bursts according to the campaign scenario definitions emitted by `run_campaign_eval.py`.

We evaluate SACD on real PRISM score streams for clean-only, sustained, burst, and low-rate scenarios. The calibrated threshold artifact records the selected hazard rate, alert probability, warmup, and false-alarm/gap metrics.

**Calibration requirement.** BOCPD requires domain-calibrated priors. With default prior ( $\mu_0 = 0.1$ ), the detector does not fire within 150 queries because the NIG posterior cannot match the actual canonical clean score mean ( $\sim 0.29$  on the ensemble logit scale; see `score_quantiles.clean.mean` in the per-seed JSONs) fast enough to yield sensitive change-point detection. For deployment,  $\mu_0$  should be set to the empirical mean of the calibration-set anomaly scores.

## D Hyperparameter Sensitivity

| Hyperparameter                       | Default            | Range tested | Sensitivity                                        |
|--------------------------------------|--------------------|--------------|----------------------------------------------------|
| $n_{\text{cal}}$ (calibration size)  | 2000               | 500–3000     | Low ( $\pm 0.03$ on L3)                            |
| $n_{\text{sub}}$ (TDA subsampling)   | 150                | 50–500       | Medium (TPR $\pm 4\%$ )                            |
| $K$ (expert clusters)                | 4                  | 2–8          | Low (Recovery $\pm 2\%$ )                          |
| $\lambda$ (L0 threshold factor)      | 0.8                | 0.6–0.95     | Medium (FPR tradeoff)                              |
| $p_{\text{alert}}$ (BOCPD threshold) | 0.60               | 0.30–0.70    | Low (latency $\pm 2$ steps)                        |
| $t_{\text{warmup}}$ (warmup steps)   | 20                 | 20–50        | Low (selected feasible cells span 20–50)           |
| $\delta$ (CUSUM drift floor)         | 1.0                | 0.00–1.0     | High at $\delta = 0$ (59.4% clean active-fraction) |
| $C^*$ (CUSUM threshold)              | 5.0                | 5.0–12.0     | Medium (long-stream durability; see §E)            |
| $k_{\text{consec}}$ (CUSUM consec.)  | 3 (C10) / 5 (C100) | 1–5          | High at $k_{\text{consec}} = 1$                    |

Table 10: Hyperparameter sensitivity summary.

$k_{\text{consec}} \times t_{\text{warmup}}$  **grid (deployed CUSUM,  $C^*=5.0$ ,  $\delta=1.0$ ).** We sweep  $k_{\text{consec}} \in \{1, 2, 3, 4, 5\}$  against  $t_{\text{warmup}} \in \{10, 20, 35, 50, 75\}$  on a 1,000-step Gaussian clean stream (mean / max active-fraction over 5 seeds) and on a sustained  $\rho=1.0$  adversarial stream (median / max time-to-detect over 5 seeds, in queries since first adversarial step).

| Cell<br>$k_{\text{consec}}$   | $t_{\text{warmup}}$ | Clean active-frac (1000 step) |       | $t$ -to-detect (queries) |     |
|-------------------------------|---------------------|-------------------------------|-------|--------------------------|-----|
|                               |                     | mean                          | max   | median                   | max |
| 1                             | any                 | 0.225                         | 0.563 | 0                        | 1   |
| 2                             | any                 | 0.144                         | 0.559 | 1                        | 2   |
| <b>3 (CIFAR-10 deployed)</b>  | 20                  | 0.112                         | 0.558 | 2                        | 3   |
| 4                             | any                 | 0.111                         | 0.557 | 3                        | 4   |
| <b>5 (CIFAR-100 deployed)</b> | 20                  | 0.111                         | 0.556 | 4                        | 5   |

Table 11:  $k_{\text{consec}}$  trades  $t$ -to-detect against clean active-fraction approximately linearly above  $k=3$ .  $t_{\text{warmup}}$  in the range 10–75 is statistically indistinguishable on these streams (the CUSUM crosses fast enough that warmup does not filter sustained adversarial signals); we therefore collapse the  $t_{\text{warmup}}$  column to “any”. The deployed CIFAR-10 cell ( $k=3$ ) is at the inflection point:  $k=2$  trades  $\approx 1$  query of latency for +3.3 pp clean active-fraction, while  $k=4$  buys  $-0.1$  pp clean for +1 query latency. CIFAR-100’s  $k=5$  trades 2 extra latency queries for negligible clean-fraction gain — it was selected because CIFAR-100’s clean score distribution has slightly heavier tails. Note that “clean active-fraction” counts every step after a single spurious trigger because the deployed monitor has no auto-reset; a deployment with periodic L0 reset would see this number drop by  $\approx 100\times$ . Source: `experiments/stress/test_10_hparam_grid.py`.

## E L0 Long-Stream Stress Test

The campaign-detection experiments in §4.7 report 0.000 clean-only L0 active-fraction on the 500-step calibration streams used by `calibrate_l0_thresholds.py`. Because deployed monitors run for longer than 500 queries, we additionally stress-test the calibrated CUSUM detector on  $\{500, 1000, 2000, 5000, 10000\}$ -step clean streams across all 5 seeds. The stream is bootstrap-resampled from  $n = 1,000$  *real* post-fusion ensemble clean scores generated end-to-end by the deployed pipeline (backbone  $\rightarrow$  TMM  $\rightarrow$  ensemble) on the VAL split (`experiments/stress/compute_real_clean_scores.py`). The empirical pool has mean 0.173 and std 2.045, matching the pkl calibration scale (`score_center= 0.281`, `score_scale= 1.639`).

Each stream is processed by a fresh CampaignMonitor loaded from `models/10_thresholds.pkl` (`detection_mode='cusum'`,  $\delta = 1.0$ ,  $t_{\text{warmup}} = 20$ ,  $k_{\text{consec}} = 3$ ), sweeping  $C^* \in \{5.0, 8.0, 10.0, 12.0\}$  (the feasible CUSUM-threshold grid cells from the calibration `top_rows`).

| $C^*$                 | 500   | 1000  | 2000  | 5000  | 10000 |
|-----------------------|-------|-------|-------|-------|-------|
| <b>5.0 (deployed)</b> | 0.000 | 0.000 | 0.019 | 0.148 | 0.274 |
| 8.0                   | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |
| 10.0                  | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |
| 12.0                  | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

Table 12: Mean clean-only L0 active-fraction across 5 seeds, computed on real post-fusion ensemble clean scores. Once L0 fires, it remains active until manually deactivated (the monitor has no auto-reset), so a single spurious trigger inflates the active-fraction over the remaining stream. The deployed  $C^* = 5.0$  holds the 0.000 claim through 1,000 steps but admits its first spurious trigger between step 1,800 and step 4,500 (1–2 of 5 seeds at the lengths shown). The alternative-feasible CUSUM thresholds from `calibrate_10_thresholds.py`’s `top_rows` ( $C^* \in \{8.0, 10.0, 12.0\}$ ) all hold 0.000 active-fraction on every seed up to 10,000 steps. Source: `experiments/stress/test_l0_real_long_clean_stream.py` operating on `experiments/stress/real_clean_scores_val_n1000.npy` (deployed-pipeline VAL-split clean scores).

**Take-away.** The deployed  $C^* = 5.0$  is selected on the calibration objective (*argmax ASR-gap subject to clean-only active-fraction  $\leq 0.01$  on a 500-step calibration stream*); it satisfies the calibration constraint at any stream length up to  $\approx 1,000$  queries and starts to admit late spurious triggers beyond  $\approx 2,000$  queries on real clean streams. For deployments where the monitor runs over thousands of queries without periodic reset, we recommend  $C^* \in [8.0, 12.0]$ , which sits at the same selected top-row ASR-gap cell (7.2 pp) in the `pkl` but is strictly more durable on long clean streams. Periodic reset (or `deactivate_10()` on return to the clean score baseline) is the orthogonal mechanism already noted in §4.7 and is required for arbitrarily long deployments regardless of  $C^*$ .

## F Per-Component Latency Breakdown

The headline latency numbers in Table 26 are end-to-end. This appendix decomposes a single PRISM forward pass into eight stages, measured on  $n=200$  clean EVAL-split images on the deployed RTX 4090 hardware via `experiments/stress/vastai_latency_breakdown.py`. The end-to-end mean here is 44.95 ms, consistent with the 39.78 ms in Table 26 (which used the same hardware on a different EVAL subsample). The breakdown identifies which channels dominate the budget — relevant for ensemble-pruning ablations and for deciding which channels are worth a closed-form differentiable surrogate in a future ensemble-complete adaptive attack (§5).

**Implications for adaptive-attack design.** The combined gradient-norm + StabilityV2 + ensemble-fusion budget ( $\sim 60\%$  of forward time) is also the portion of the pipeline that the canonical adaptive attack (§4.4) does not target, because StabilityV2’s  $M=5$  jitter-stability statistic has no closed-form differentiable surrogate, and gradient-norm is itself a gradient quantity. A genuine ensemble-complete BPDA must either (i) construct a surrogate (e.g., end-to-end backprop through  $M$  noise draws with straight-through estimator on the argmax), or (ii) accept the defender’s untargeted budget as an attack-cost asymmetry.

## G High-Confidence CW Sensitivity ( $\kappa$ Sweep)

The canonical CW evaluation in Table 2 uses  $\kappa=0$ . Carlini Carlini and Wagner [2017] recommends a confidence-margin sweep for robust detector evaluation, so we additionally measure TPR at  $\kappa \in \{0, 10, 20\}$  on  $n=1000$  EVAL-split inputs per seed across all 5 fixed seeds (paper-grade protocol matching Table 2), using the native PyTorch CW implementation (`src/attacks/cw_torch.py`) at the *paper-canonical* attack budget `max_iter=100`, `bss=9`. All three configurations achieve  $\geq 99.8\%$  attack-success against the unwrapped backbone with mean  $L_2 \leq 0.151$ .

| Stage                                  | mean (ms)    | p50 (ms) | p95 (ms) | % of total   |
|----------------------------------------|--------------|----------|----------|--------------|
| Backbone forward (logits)              | 1.98         | 1.98     | 2.03     | 4.4%         |
| TAMM activation extraction             | 2.00         | 2.00     | 2.05     | 4.5%         |
| TDA persistence diagrams (3 layers)    | <b>14.24</b> | 14.06    | 16.29    | <b>31.7%</b> |
| Input gradient norm (autograd back)    | 7.07         | 6.35     | 6.72     | 15.7%        |
| Logit-profile features                 | 0.11         | 0.10     | 0.15     | 0.2%         |
| StabilityV2 ( $M=5$ jittered fwd)      | 8.81         | 8.77     | 8.89     | 19.6%        |
| DCT image prep                         | 0.01         | 0.01     | 0.01     | 0.0%         |
| Ensemble logistic + Wasserstein fusion | 10.72        | 10.25    | 14.13    | 23.8%        |
| <b>Total</b>                           | <b>44.95</b> | 43.72    | 48.38    | 100.0%       |

Table 13: Per-stage latency on RTX 4090 ( $n=200$  clean EVAL samples, single seed). TDA persistence-diagram computation is the largest single contributor at 31.7% (14.24 ms), consistent with the §5 limitation note that persistent homology dominates the inference budget. Ensemble fusion (23.8%), StabilityV2 (19.6%), and gradient-norm (15.7%) follow in order. We additionally ran the same breakdown on a GTX 1060 6GB and observed StabilityV2 dominating (30.3% of a total 71.18 ms) while TDA dropped to 17.3%: TDA is CPU-bound (gudhi auction algorithm on the persistence-diagram side), so it benefits little from a faster GPU, whereas StabilityV2’s  $M=5$  forward passes are GPU-bound and scale proportionally with backbone throughput. The relative ordering in this table is therefore RTX-4090-specific; consumer GPUs will see StabilityV2 climb. Source: `experiments/stress/vastai_latency_breakdown.json`.

| $\kappa$      | Attack success | mean $L_2$ | TPR @ L1 [95% CI]            | TPR @ L2 | TPR @ L3 |
|---------------|----------------|------------|------------------------------|----------|----------|
| 0 (canonical) | 0.998          | 0.128      | <b>0.9486</b> [0.942, 0.954] | 0.805    | 0.339    |
| 10            | 0.998          | 0.149      | <b>0.7762</b> [0.764, 0.788] | 0.565    | 0.194    |
| 20            | 0.999          | 0.151      | <b>0.7906</b> [0.779, 0.802] | 0.578    | 0.198    |

Table 14: CW confidence-margin sensitivity at paper-canonical attack budget ( $\text{max\_iter}=100$ ,  $\text{bss}=9$ ), 5 seeds  $\times n=1000 = 5,000$  EVAL adversaries per  $\kappa$  on RTX 4090 with the deployed scorer. Pooled clean FPR on the same EVAL pool is  $L1 = 0.078$ ,  $L2 = 0.017$ ,  $L3 = 0.002$ . The  $\kappa=0$  headline (0.9486) lands within 1 pp of the main-table value (0.938, Table 2), confirming the canonical result is reproducible under the stronger attack budget. Higher-confidence CW reduces TPR at L1 by  $\approx 17$  pp from  $\kappa=0$  to  $\kappa=10$  (0.776), *plateauing* at  $\kappa=20$  (0.791, statistically tied with  $\kappa=10$ ) rather than further collapsing. The detector still discriminates adversarial from clean (0.78 vs 0.08 at L1, a  $10\times$  ratio) at the strongest tested margin. The drop is consistent with the ensemble being trained on FGSM/PGD/Square only (CW is *not* in the training mix; see Appendix L): the canonical  $\kappa=0$  CW signature is caught via transfer from the gradient-following statistics the ensemble has learned, whereas higher-confidence ( $\kappa>0$ ) CW drifts off that transferable signature. Per-attack training augmentation with  $\kappa>0$  (or CW) samples is a natural robustness lever for future work. Source: `experiments/stress/vastai_stronger_cw_canonical.json`.

## H AutoAttack- $L_\infty$ Evaluation

AutoAttack Croce and Hein [2020] is the de-facto standard adversarial-robustness benchmark, combining four parameter-free attacks (APGD-CE, APGD-DLR, FAB-T, and Square) into a strong, reproducible ensemble. We evaluate the deployed PRISM detector at the canonical  $\varepsilon_\infty=8/255$  budget on  $n=1000$  EVAL-split inputs per seed across all 5 fixed seeds (paper-grade protocol matching Table 2). Each sub-attack is run independently so per-attack TPR can be reported; Square is constrained to  $n_{\text{queries}}=1000$  for tractability.

## I Extended L3 Recovery (Multi-Attack)

The main recovery table (Table 9) reports L3 recovery only on PGD adversaries. We additionally measure recovery on FGSM, Square, and CW-L2 on  $n=1000$  EVAL-split inputs per seed across all 5 fixed seeds (paper-grade protocol matching Table 9) to characterise how the TAMSH router generalises across attack families. ART’s FGSM ( $\varepsilon=8/255$ ) and SquareAttack ( $\text{max\_iter}=5000$ , 1 restart) generate the perturbations; CW uses the same canonical  $\kappa=0$  configuration as Table 14.

| Sub-attack | Attack success | mean $L_\infty$ | TPR @ L1 [95% CI]            | TPR @ L2 | TPR @ L3      |
|------------|----------------|-----------------|------------------------------|----------|---------------|
| APGD-CE    | 0.999          | 0.0299          | <b>0.9708</b> [0.966, 0.975] | 0.9608   | <b>0.9500</b> |
| APGD-DLR   | 1.000          | 0.0299          | <b>0.9112</b> [0.903, 0.919] | 0.7782   | 0.3156        |
| FAB-T      | 0.971          | 0.0038          | <b>0.9314</b> [0.924, 0.938] | 0.8338   | 0.4494        |
| Square     | 0.830          | 0.0246          | <b>0.7062</b> [0.693, 0.719] | 0.5080   | 0.1590        |

Table 15: AutoAttack- $L_\infty$  ( $\varepsilon=8/255$ ) per-sub-attack TPR on 5 seeds  $\times n=1000 = 5,000$  EVAL adversarials per attack (RTX 4090, deployed scorer). Pooled clean FPR on the same EVAL pool is L1 = 0.078, L2 = 0.017, L3 = 0.002. Three observations: (i) The gradient-based AA components (APGD-CE / APGD-DLR / FAB-T) are all detected at  $\geq 0.91$  TPR @ L1, with APGD-CE persisting at 0.95 TPR even at the high-confidence L3 threshold — the detector treats AA-generated perturbations as far higher-confidence anomalies than PGD at the same budget (cf. Table 2 PGD L1 = 0.987, but with much shallower L3 tail). (ii) FAB-T converges with  $L_\infty \approx 0.0038$  ( $\sim 8\times$  under the budget) yet is still detected at 0.93 TPR @ L1; the adversarial signature is visible long before the perturbation budget is exhausted. (iii) The black-box Square component achieves the lowest TPR (0.706 @ L1) at the highest attack-failure rate (17%). This does *not* exceed the strongest Square-specific baseline (Energy at 0.974 Square TPR, Table 7); but Energy buys that Square number at the cost of collapsing to 0.001 on PGD, whereas PRISM holds  $\geq 0.71$  TPR on *every* AutoAttack sub-attack and  $\geq 0.88$  across the four headline attacks — Energy has no robust mean-TPR floor across attacks. Source: experiments/stress/vastai\_autoattack.json.

| Attack | $n_{L3}$ pool | Trigger rate | Passthrough | Uniform router         | Topology router | Oracle |
|--------|---------------|--------------|-------------|------------------------|-----------------|--------|
| FGSM   | 783           | 0.157        | 0.346       | 0.420 (+7.4pp)         | 0.377 (+3.1pp)  | 0.670  |
| Square | 1211          | 0.242        | 0.014       | <b>0.486</b> (+47.2pp) | 0.449 (+43.5pp) | 0.893  |
| CW     | 1254          | 0.251        | 0.017       | <b>0.487</b> (+47.0pp) | 0.407 (+39.0pp) | 0.857  |

Table 16: L3-recovery on FGSM/Square/CW adversarials, pooled across 5 seeds  $\times n=1000 = 5,000$  EVAL adversarials per attack (RTX 4090, deployed scorer). “Topology router” is the combined  $H_0+H_1$  Wasserstein gate, matching §4.8; Wilson 95% CIs on the topology row are FGSM [0.343, 0.411], Square [0.421, 0.477], CW [0.380, 0.434]. Four points stand out: (i) On **CW**, the topology router recovers 40.7% of L3-rejected adversarials from a 1.7% passthrough baseline (+39.0 pp), nearly  $4\times$  the paper’s PGD-only +10.6 pp main result. (ii) On **Square**, the topology router lifts recovery to 44.9% (+43.5 pp), confirming the C4 contribution transfers beyond white-box gradient attacks. (iii) On **FGSM**, the topology router gives only +3.1 pp because the unrouted backbone already recovers 34.6% of L3 cases (FGSM at  $\varepsilon=8/255$  is a low-budget single-step perturbation that does not always flip the argmax even when L3 triggers). Force-routing to the PGD specialist (Table 9) is catastrophic on every cross-attack (0.10, 0.04, 0.05 recovery respectively), confirming that the prior-augmented variant is a PGD-specific upper-bound, not a deployable cross-attack policy. (iv) The **uniform-router control** consistently matches or beats the topology router across all three attacks; for the deployed expert pool (trained on clean/FGSM/PGD/mixed) the topology signal does *not* add over uniform  $1/K$  mixing on non-PGD attacks. The PGD-only paper claim that uniform is degenerate to passthrough (Table 9, tamsh\_uniform paired gap 0 pp) is therefore attack-specific and does not extend to FGSM/Square/CW; on these attacks uniform mixing of the attack-specialist experts beats argmax of the unrouted backbone. Oracle (per-input best-of-4) ceilings of 0.89/0.86 on Square/CW indicate substantial remaining headroom for a stronger learned router. Source: experiments/stress/vastai\_recovery\_multi.json.

**Take-away.** TAMSH’s C4 contribution is *attack-family-dependent* and the ranking of router strategies shifts across attacks. CW and Square benefit substantially from *any* mixture of the deployed experts (topology or uniform); FGSM does not, and the paper’s “uniform = degenerate” claim from Table 9 does not transfer beyond PGD. Future work on TAMSH should treat per-attack-family recovery and uniform-vs-topology comparison as the primary metrics, train an FGSM specialist (the current expert pool under-serves single-step perturbations), and characterise the gap between uniform mixing and topology routing across the attack matrix.

## J Per-Attack Score Histograms (Local Reproduction)

Figure 7 shows clean vs adversarial score distributions for FGSM, Square, and CW on a single-seed VAL-split run ( $n=80$ ), with the calibrated L1/L2/L3 thresholds overlaid. All three adversarial distributions are clearly separated from the clean distribution at L1; the long tail of CW into the L3 region is the dominant driver of the high CW recovery gap in Table 16.




Figure 7: Per-attack PRISM score distributions on VAL-split (single-seed,  $n=80$ , local reproduction). Clean (blue) vs adversarial (orange/green/red) histograms with vertical dashed lines at L1/L2/L3 conformal thresholds. Clean per-tier coverage is  $L1 = 0.10$ ,  $L2 = 0.02$ ,  $L3 = 0.00$ , consistent with the calibrated  $\alpha$  values. Source: `experiments/stress/test_score_histograms.py`.

| Set    | mean | std  | p50  | p95  | TPR @ L1 | TPR @ L2 |
|--------|------|------|------|------|----------|----------|
| clean  | 0.27 | 2.10 | 0.28 | 3.43 | 0.100    | 0.020    |
| FGSM   | 5.36 | 1.86 | 5.31 | 8.63 | 0.925    | 0.638    |
| Square | 5.60 | 2.12 | 5.97 | 8.72 | 0.863    | 0.638    |
| CW     | 5.92 | 2.34 | 6.03 | 9.12 | 0.925    | 0.700    |

Table 17: Per-set score quantiles on the local single-seed VAL run. Adversarial means cluster around 5.4–5.9, well above the L1 threshold 2.811 but only partially above the L3 threshold 7.108.

## K Appendix-Only: SACD ASR Gap (Demoted Metric)

The naive “ASR gap (L0 off vs. L0 on)” metric at  $\rho=1.0$  is reported here for completeness; the headline campaign metric is time-to-detect (§4.7, Table 8).

## L Appendix-Only: Weakened-CW Detector-Fidelity Ablation

A 5-seed local rerun with weakened CW parameters ( $\text{max\_iter} = 5$ ,  $\text{bss} = 5$ , vs. the canonical  $\text{max\_iter} = 100$ ,  $\text{bss} = 9$  in Table 2) shows that the ensemble detector’s CW TPR depends on attack-signature fidelity, not on perturbation magnitude:

## M Appendix-Only: TAMSH Router Diagnosis (F4 Fix)

The pre-fix TAMSH router used  $H_1$ -only Wasserstein distance for expert gating. Diagnostic audit (seed 42,  $n = 200$  PGD inputs) showed:

| Scenario                 | ASR (L0 off)      | ASR (L0 on)       | $\Delta$ (pp, paired) |
|--------------------------|-------------------|-------------------|-----------------------|
| clean_only               | 0.000 $\pm$ 0.000 | 0.000 $\pm$ 0.000 | +0.00 $\pm$ 0.00      |
| sustained_ $\rho$ =0.5   | 0.015 $\pm$ 0.014 | 0.011 $\pm$ 0.015 | +0.39 $\pm$ 0.88      |
| sustained_ $\rho$ =0.8   | 0.012 $\pm$ 0.009 | 0.009 $\pm$ 0.010 | +0.24 $\pm$ 0.54      |
| sustained_ $\rho$ =1.0   | 0.011 $\pm$ 0.007 | 0.009 $\pm$ 0.009 | +0.20 $\pm$ 0.45      |
| burst                    | 0.010 $\pm$ 0.016 | 0.010 $\pm$ 0.016 | +0.00 $\pm$ 0.00      |
| low_rate ( $\rho$ =0.05) | 0.004 $\pm$ 0.009 | 0.000 $\pm$ 0.000 | +0.41 $\pm$ 0.91      |

Table 18: Campaign-stream ASR gap (5-seed mean  $\pm$  std). The ceiling is structurally bounded: the upstream L3 conformal threshold rejects  $\sim 99\%$  of  $\rho=1.0$  adversarials within 22 queries, leaving no headroom for an L0-attributable ASR gap (all gaps  $\leq 0.41$  pp). Time-to-detect, not ASR gap, is the operationally meaningful campaign metric (Table 8).

| CW configuration                    | Mean L2 distortion | Attack success | Detector TPR                        |
|-------------------------------------|--------------------|----------------|-------------------------------------|
| canonical (max_iter = 100, bss = 9) | 0.128              | 1.000          | <b>0.938 <math>\pm</math> 0.008</b> |
| weakened (max_iter = 5, bss = 5)    | 0.193              | 0.896          | 0.645 $\pm$ 0.013                   |

Table 19: Stronger CW with smaller distortion is *easier* to detect than weakened CW with larger distortion. Converged CW produces a clean “minimum-distortion-misclassification” signature in the logit-profile and DCT channels that the ensemble has learned (the ensemble is trained on FGSM/PGD/Square; CW signature transfers because optimisation-based attacks share these gradient-following statistics with PGD). Non-converged CW does not produce this signature and falls in a detector blind spot.

1.  $H_1$  medoid diagrams contain only 1–3 points per expert on layer4 of a CIFAR-native ResNet-18.
2. 22.7% of L3-rejected inputs have *empty*  $H_1$  diagrams.
3. The empty-diagram fallback in the Wasserstein implementation returns  $\sum |b - d|$ , which evaluates to the same tiny value for all four experts. `argmin` then defaults to index 0.
4. Net effect: expert 0 (clean-trained) wins 94% of routes ( $n = 163$ , picks=[154, 0, 8, 1]); routing entropy = 0.336/2.000 bits.
5. Force-routing the same inputs to each expert in isolation shows expert 2 (PGD-trained) achieves 35% recovery accuracy vs.  $\sim 5\%$  for experts 0/1/3. The experts are individually capable; only the routing signal is broken.

The F4 fix adds an  $H_0$  term to the routing metric (Eq. 5):  $H_0$  medoid diagrams have  $\sim 16$  points (one per connected component at large filtration radius) and are sufficiently dense for Wasserstein gating to be informative. On the 5-seed L3-rejected pool, per-input oracle (best-of- $K$ ) ceiling is 39.9% recovery; the combined-metric router achieves 13.9% (35% of oracle, +10.6 pp over passthrough); always-route-to-expert-2 achieves 29.3% (73% of oracle, +26.0 pp over passthrough, clearing the +15 pp C4 gate).

## N CIFAR-100 Results

We report CIFAR-100 results following the same protocol as §4 (5 seeds,  $n = 1000$  per seed; CIFAR-native ResNet-18 trained from scratch; conformal calibration on a held-out clean split with the same three FPR tiers L1/L2/L3 = 0.10, 0.03, 0.005). Headline detection numbers are in Table 3; this appendix collects the per-component artifacts.

**High-confidence CW sensitivity (CIFAR-100).** We additionally measure CW TPR at  $\kappa \in \{0, 10, 20\}$  on the CIFAR-100/ResNet-18 setting, mirroring the CIFAR-10 protocol of Table 14. Five seeds  $\times n=1000$  EVAL-split adversarials per  $\kappa$  with the paper-canonical CW configuration (max\_iter=100, bss=9) yields the numbers in Table 20. Two observations: (i) the  $\kappa=0$  headline (0.7212) is  $\sim 23$  pp below the CIFAR-10 canonical CW TPR (0.9486, Table 14), consistent with the CIFAR-100 result being harder for optimisation-based attacks across the board (Table 3); (ii) the  $\kappa=0 \rightarrow \kappa=10$  drop is  $\sim 10$  pp on CIFAR-100 vs  $\sim 17$  pp on CIFAR-10, and  $\kappa=10$  and  $\kappa=20$  are statistically tied on both datasets — the detector saturates rather than collapses under high-confidence



CW. At the strongest tested margin ( $\kappa=20$ ) the detector still produces a  $8.7\times$  discrimination ratio between adversarial (0.612) and clean (0.073) at L1 on CIFAR-100.

| $\kappa$      | Attack success | mean $L_2$ | TPR @ L1 [95% CI]            | TPR @ L2 | TPR @ L3 |
|---------------|----------------|------------|------------------------------|----------|----------|
| 0 (canonical) | 0.994          | 0.079      | <b>0.7212</b> [0.709, 0.733] | 0.3908   | 0.0866   |
| 10            | 0.992          | 0.085      | <b>0.6098</b> [0.596, 0.623] | 0.3104   | 0.0716   |
| 20            | 0.993          | 0.085      | <b>0.6124</b> [0.599, 0.626] | 0.3106   | 0.0702   |

Table 20: CIFAR-100 CW confidence-margin sensitivity at paper-canonical attack budget (max\_iter=100, bss=9), 5 seeds  $\times n=1000 = 5,000$  EVAL adversarials per  $\kappa$  on RTX 4090 with the deployed CIFAR-100 scorer. Pooled clean FPR on the same EVAL pool is L1 = 0.073, L2 = 0.025, L3 = 0.002. Source: experiments/stress/vastai\_stronger\_cw\_canonical\_c100.json.

**Adaptive PGD ( $\lambda$  sweep).** Table 21 reports adaptive PGD (BPDA, activation-matching + DCT-energy matching via -through-scorer) with 100 steps and 10 restarts across the  $\lambda$  sweep used in §4.4. CIFAR-100 PRISM is robust to the adaptive attack: TPR holds in  $[0.92, 0.95]$  across  $\lambda \in [0, 10]$ , in marked contrast to the CIFAR-10 result (Table 5), where the same attack collapses TPR from 0.832 at  $\lambda=0$  to 0.429 at  $\lambda=10$ .

**Why CIFAR-100 resists the adaptive attack while CIFAR-10 collapses.** The attack’s suppression term covers two ensemble channels: per-layer activation matching (TAMM target) and DCT high-frequency energy matching. It does *not* suppress the StabilityV2, gradient-norm, or softmax-entropy channels, which the attack would need separate loss terms (and, in StabilityV2’s case, a differentiable surrogate of the jitter-stability statistic) to attack jointly. On CIFAR-10, the channel-mask ablation (Table 6) shows that TDA and DCT together carry the dominant FGSM/PGD discriminative weight, so suppressing them removes most of the detector’s signal and TPR falls. On CIFAR-100, the ablation in Table 22 and the cross-setting analysis in §4.3 identify StabilityV2 as the dominant non-PGD channel (No-StabilityV2 collapses PGD TPR to 0.423, Square to 0.279, CW to 0.205, and FGSM to 0.562, with FPR dropping to 0.019). Because the adaptive loss does not attack StabilityV2, the channel continues to fire on CIFAR-100 adversarials regardless of  $\lambda$ , and the L3-tier trigger rate stays near 88% even at  $\lambda=10$ . The robust CIFAR-100 result is therefore not evidence that PRISM is universally adaptive-secure; it is evidence that the attack covers the wrong channels for this dataset. A genuinely ensemble-complete adaptive attack — adding differentiable surrogates for the StabilityV2, gradient-norm, and softmax-entropy losses — is the required next step (see §5); the  $\lambda$ -sweep results bound the partial-coverage attacker.

The corresponding WRN-28-10 adaptive-PGD evaluation is reported in §O (Table 27): WRN behaves identically to ResNet-18 on CIFAR-10 under the same attack, so the architecture-agnostic interpretation is that the channel-coverage analysis above governs both settings.

| $\lambda$ | TPR   | FPR   | $n_{\text{adv}}$ |
|-----------|-------|-------|------------------|
| 0.0       | 0.950 | 0.073 | 5000             |
| 0.5       | 0.939 | 0.073 | 5000             |
| 1.0       | 0.933 | 0.073 | 5000             |
| 2.0       | 0.922 | 0.073 | 5000             |
| 5.0       | 0.925 | 0.073 | 5000             |
| 10.0      | 0.925 | 0.073 | 5000             |

Table 21: Adaptive PGD on CIFAR-100 (ResNet-18), 5 seeds pooled.

**Channel-mask ablation (with FPR disclosure).** The CIFAR-100 ablation in Table 22 mirrors the CIFAR-10 methodology: each row removes one feature family while reusing the Full-PRISM threshold. Unlike CIFAR-10 (where No-LogitProfile inflates FPR to 0.938), on CIFAR-100 only *one* configuration inflates FPR rather than improving discrimination: No-GradNorm (FPR 0.658, i.e. flags two-thirds of clean inputs) appears to “improve” TPR only because its calibrated threshold has collapsed. The other masked rows hold FPR at or below the 0.10 design target (No-LogitProfile 0.077, No-DCT 0.058, No-StabilityV2 0.019) and mostly *lose* TPR on removal. We surface FPR directly via the explicit column so the table cannot be misread as an argument for dropping any single

channel. The legitimate signal is that StabilityV2 (PGD/Square/CW TPR collapses on its removal), TDA (FGSM TPR collapses without it), DCT, and LogitProfile each contribute non-redundant discriminative information at the calibrated FPR.

| Configuration              | FGSM TPR | PGD TPR | Square TPR | CW TPR | FPR          |
|----------------------------|----------|---------|------------|--------|--------------|
| Full PRISM (ablation cal.) | 0.832    | 0.980   | 0.673      | 0.679  | <b>0.073</b> |
| No-LogitProfile            | 0.694    | 0.974   | 0.633      | 0.541  | 0.077        |
| No-StabilityV2             | 0.562    | 0.423   | 0.279      | 0.205  | 0.019        |
| No-DCT                     | 0.730    | 0.974   | 0.581      | 0.640  | 0.058        |
| No-GradNorm                | 0.987    | 0.998   | 0.996      | 0.998  | <b>0.658</b> |
| Ensemble-no-TDA            | 0.544    | 0.973   | 0.500      | 0.710  | 0.067        |
| TDA only                   | 0.020    | 0.934   | 0.057      | 0.058  | 0.086        |

Table 22: CIFAR-100 channel-mask ablation (5 seeds,  $n = 1000$ , per-attack TPR + FPR). Rows whose FPR exceeds the 0.10 design target indicate threshold collapse, not improved detection; they are reported here for symmetry with Table 6 only. **Baseline row.** “Full PRISM (ablation cal.)” is the *ablation-internal* baseline: the ablation pipeline refits one split-conformal threshold shared by every channel-mask arm, so each masked row is measured against a common operating point. The canonical Full PRISM numbers are the CIFAR-100 row of Table 3 (FGSM 0.835, PGD 0.984, Square 0.672, CW 0.687, FPR 0.077), from the detection pipeline’s canonical calibrator on the *same* eval split [8000, 10000], 5 seeds,  $n=1000$ . The two independent calibration fits give operating points that differ by 0.1–0.8 pp TPR (mean 0.4 pp; largest on CW, 0.679 vs 0.687) and 0.4 pp FPR (0.073 vs 0.077); every value here lies within the Wilson 95% CIs of Table 3, which is the row to cite as canonical.

**Campaign detection (CIFAR-100).** Table 23 reports CIFAR-100 ASR with the L0 monitor off vs. on across the same six stream scenarios. As in the CIFAR-10 case, the ASR gap is structurally bounded (§4.7): the upstream L1/L3 conformal already crushes campaign ASR to  $\leq 4.3\%$ , leaving at most 1.50 pp of additional headroom for L0 across any scenario. The headline metric (time-to-detect) is the operationally meaningful one; gap is reported for completeness and is consistent with the demotion clause pre-registered in plan P0.4.

| Scenario                 | ASR (L0 off) | ASR (L0 on) | $\Delta$ (pp) |
|--------------------------|--------------|-------------|---------------|
| clean_only               | 0.000        | 0.000       | +0.00         |
| sustained_ $\rho=0.5$    | 0.036        | 0.026       | +1.07         |
| sustained_ $\rho=0.8$    | 0.032        | 0.019       | +1.29         |
| sustained_ $\rho=1.0$    | 0.031        | 0.018       | +1.30         |
| burst                    | 0.043        | 0.028       | +1.50         |
| low_rate ( $\rho=0.05$ ) | 0.013        | 0.004       | +0.85         |

Table 23: CIFAR-100 campaign-stream ASR with L0 off vs. on (5 seeds, post-fix). Clean-only false alarm rate is 0.000; ASR-reduction gaps are structurally bounded by L1 conformal action.

**L3 recovery (CIFAR-100).** Recovery accuracy on CIFAR-100 (Table 24) is reported for completeness; all routing strategies sit within 1.4 pp of passthrough and well below the +15 pp gate target. The 100-class setting collapses both the passthrough baseline (chance accuracy 0.01) and the recovery ceiling: the CIFAR-10 tamsh\_force upper-bound (+26 pp; Table 9) does not transfer to CIFAR-100, indicating that the PGD-trained expert’s specialization no longer corresponds to a recoverable label cluster when the class space is  $10\times$  larger. We accept this miss per the pre-registered demotion clause in plan P0.5: TAMSH (C4) is reported as an appendix-only secondary result on CIFAR-100.

**L3 recovery on FGSM/Square/CW (CIFAR-100).** The PGD-only collapse above does *not* extend to the rest of the attack mix. Re-running the same recovery harness with FGSM, Square, and CW-L2 adversarials on  $n=1000$  EVAL-split inputs per seed across all 5 fixed seeds reveals a strongly attack-conditional recovery profile (Table 25). The topology router recovers +42.4 pp over passthrough on CW and +31.4 pp on Square — exceeding the C4 +15 pp gate by a wide margin — while remaining

| Strategy      | Recovery Acc. | 95% CI         | Availability | $\Delta$ vs passthrough |
|---------------|---------------|----------------|--------------|-------------------------|
| reject        | 0.000         | [0.000, 0.001] | 0.000        | —                       |
| passthrough   | 0.056         | [0.050, 0.063] | 1.000        | —                       |
| tamsh         | 0.065         | [0.058, 0.072] | 1.000        | +0.91 pp                |
| tamsh_uniform | 0.057         | [0.051, 0.064] | 1.000        | +0.13 pp                |
| tamsh_force   | 0.070         | [0.063, 0.077] | 1.000        | +1.36 pp                |

Table 24: CIFAR-100 L3 recovery on PGD adversarials (5 seeds pooled). All TAMSH strategies sit within 1.4 pp of passthrough; the +15 pp C4 gate is missed and the component is reported as appendix-only.

neutral or marginally negative on FGSM. The uniform  $1/K$  control is statistically tied with the topology router on Square/CW (same pattern as the PGD-on-CIFAR-10 result in Table 9), confirming that recovery on optimization- and query-based attacks is driven by expert capacity at the calibrated FPR, with the topology gate not adding measurable signal above  $1/K$ . The force-PGD specialist that delivers +26 pp on CIFAR-10/PGD collapses to  $\leq +7.1$  pp on CIFAR-100 for every attack family, confirming that the PGD-trained expert’s specialization no longer corresponds to a recoverable label cluster in the 100-class setting. The FGSM column tracks the CIFAR-10 FGSM result (Table 16) qualitatively (small, mixed-sign deltas) but the per-seed signal is too noisy at  $n_{L3} \approx 185$  per seed to claim a positive recovery; we report the negative outcome honestly.

| Attack | $n_{L3}$ pool | Trigger rate | Passthrough | Uniform router         | Topology router        | Oracle |
|--------|---------------|--------------|-------------|------------------------|------------------------|--------|
| FGSM   | 924           | 0.185        | 0.121       | 0.142 (+2.1pp)         | 0.102 (−1.9pp)         | 0.306  |
| Square | 449           | 0.090        | 0.020       | <b>0.354</b> (+33.4pp) | <b>0.334</b> (+31.4pp) | 0.626  |
| CW     | 394           | 0.079        | 0.015       | <b>0.447</b> (+43.1pp) | <b>0.439</b> (+42.4pp) | 0.744  |

Table 25: CIFAR-100 L3-recovery on FGSM/Square/CW adversarials, pooled across 5 seeds  $\times n=1000 = 5,000$  EVAL adversarials per attack (RTX 4090, deployed C100 scorer,  $L_3$  threshold = 5.584). Wilson 95% CIs on the topology row are FGSM [0.084, 0.123], Square [0.292, 0.379], CW [0.391, 0.488]. The +42.4 pp CW and +31.4 pp Square recovery decisively exceed the +15 pp C4 gate, so the C4 demotion clause in plan P0.5 applies only to the PGD/FGSM sub-cases on CIFAR-100. Source: experiments/stress/vastai\_recovery\_multi\_c100.json.

**Latency.** Single-sample end-to-end PRISM detection latency (Table 26) is well below the 100 ms real-time target on both datasets/backbones. Maximum-latency outliers reflect first-batch CUDA kernel compilation; steady-state p95 is the deployable figure.

| Dataset   | Backbone  | Mean (ms) | p50 (ms) | p95 (ms) | Max (ms) | Target (ms) |
|-----------|-----------|-----------|----------|----------|----------|-------------|
| CIFAR-10  | ResNet-18 | 39.78     | 35.92    | 41.26    | 724.31   | 100         |
| CIFAR-100 | ResNet-18 | 50.35     | 44.94    | 55.90    | 841.60   | 100         |

Table 26: End-to-end PRISM detection latency (RTX 4090,  $n = 200$ ).

## O WRN-28-10 Adaptive PGD

The headline adaptive-PGD result (Table 5, §4.4) uses CIFAR-10/ResNet-18. To check whether the  $\lambda$ -sweep degradation is architecture-specific, we re-run the same attack on WRN-28-10/CIFAR-10 with identical hyperparameters ( $n=1000$  EVAL adversarials, 100 PGD steps, 10 restarts,  $\varepsilon=8/255$ , BPDA via `-through-scorer`). Results in Table 27.

**Take-away.** The adaptive-PGD vulnerability is dataset-specific (CIFAR-10), not architecture-specific. WRN-28-10 inherits the same ensemble-channel-coverage gap that the CIFAR-10/ResNet-18 adaptive attack exploits (§4.4: TDA + DCT carry dominant discriminative weight on CIFAR-10; suppressing them via `-through-scorer` removes most of the detector’s signal). The CIFAR-100 robustness gap (Table 21, TPR  $\in [0.92, 0.95]$  across the same  $\lambda$  sweep) confirms the same mechanism

| $\lambda$ | WRN-28-10 TPR [95% CI]       | FPR   | ResNet-18 TPR | $\Delta$ |
|-----------|------------------------------|-------|---------------|----------|
| 0         | <b>0.8420</b> [0.818, 0.863] | 0.081 | 0.832         | +1.0 pp  |
| 0.5       | <b>0.7990</b> [0.773, 0.823] | 0.081 | 0.796         | +0.3 pp  |
| 1         | <b>0.7450</b> [0.717, 0.771] | 0.081 | 0.753         | −0.8 pp  |
| 2         | <b>0.6640</b> [0.634, 0.693] | 0.081 | 0.660         | +0.4 pp  |
| 5         | <b>0.4880</b> [0.457, 0.519] | 0.081 | 0.490         | −0.2 pp  |
| 10        | <b>0.4260</b> [0.396, 0.457] | 0.081 | 0.429         | −0.3 pp  |

Table 27: WRN-28-10/CIFAR-10 adaptive PGD ( $n=1000$ , seed 42, 100 steps  $\times 10$  restarts,  $\varepsilon=8/255$ , BPDA through the scorer). FPR holds at the calibrated 0.081 across all  $\lambda$  (conformal coverage preserved). The ResNet-18 column reproduces the 5-seed pool from Table 5 ( $n=5,000$  each, FPR 0.075); WRN here is a single-seed  $n=1000$  rerun, so CIs are wider, but the two architectures agree within  $\pm 1.0$  pp at every  $\lambda$  on a  $0.84 \rightarrow 0.43$  collapse curve. The architecture-agnostic interpretation is that adaptive-PGD vulnerability on CIFAR-10 is governed by the ensemble-channel coverage of the attack, not by the backbone choice. Source: experiments/stress/wrn/vastai\_wrn\_adaptive\_pgd\*.json (consolidated in vastai\_wrn\_adaptive\_pgd\_merged.json).

in reverse: on CIFAR-100, StabilityV2 is the dominant signal and is not covered by the adaptive loss. An ensemble-complete adaptive attack (adding differentiable surrogates for StabilityV2, gradient-norm, and softmax-entropy) is required to re-evaluate WRN-28-10 and CIFAR-100; see §5.

## P AutoAttack $L_\infty$ Epsilon Sensitivity

The canonical AA evaluation (§H, Table 15) uses  $\varepsilon=8/255$ . To measure how PRISM’s detection rate varies with the adversary’s budget, we repeat the same protocol (5 seeds  $\times n=1000$  EVAL-split, sub-attacks APGD-CE/APGD-DLR/FAB-T/Square,  $n_{\text{queries}}=1000$  for Square) at two larger budgets  $\varepsilon \in \{12/255, 16/255\}$  on the deployed CIFAR-10/ResNet-18 scorer.

| Sub-attack | TPR @ L1 |        |        | TPR @ L2 |        |        | TPR @ L3 |        |        |
|------------|----------|--------|--------|----------|--------|--------|----------|--------|--------|
|            | 8/255    | 12/255 | 16/255 | 8/255    | 12/255 | 16/255 | 8/255    | 12/255 | 16/255 |
| APGD-CE    | 0.9708   | 0.9708 | 0.9708 | 0.9608   | 0.9620 | 0.9620 | 0.9500   | 0.9556 | 0.9560 |
| APGD-DLR   | 0.9112   | 0.9560 | 0.9646 | 0.7782   | 0.8960 | 0.9312 | 0.3156   | 0.5780 | 0.7834 |
| FAB-T      | 0.9314   | 0.9290 | 0.9308 | 0.8338   | 0.8326 | 0.8308 | 0.4494   | 0.4482 | 0.4460 |
| Square     | 0.7062   | 0.8162 | 0.8714 | 0.5080   | 0.5766 | 0.6380 | 0.1590   | 0.1578 | 0.1898 |

Table 28: AutoAttack  $L_\infty$  TPR vs.  $\varepsilon$  budget, 5 seeds  $\times n=1000 = 5,000$  EVAL adversarials per (sub-attack,  $\varepsilon$ ) cell (RTX 4090). The  $\varepsilon=8/255$  column reproduces Table 15 for direct comparison;  $\varepsilon \in \{12/255, 16/255\}$  come from vastai\_autoattack\_eps\_sweep.json. Pooled clean FPR on the same EVAL pool is L1 = 0.078, L2 = 0.017, L3 = 0.002. Four observations: (i) APGD-CE detection is exactly flat at L1 (0.9708 identically across all three budgets) with mild +0.6 pp growth at L3—larger perturbations do not improve evasion against the topological signal. (ii) APGD-DLR exhibits the strongest budget sensitivity: L3 TPR climbs +47 pp from  $\varepsilon=8/255$  (0.316) to  $\varepsilon=16/255$  (0.783), confirming that the L3 conformal tier reliably catches the higher-distortion adversarials the DLR loss is willing to produce. (iii) Square (black-box) TPR increases monotonically at L1 and L2 (+11 pp / +13 pp from 8/255 to 16/255); larger random perturbations produce stronger activation anomalies without any gradient information. (iv) FAB-T converges far below the nominal budget (mean  $L_\infty \approx 0.004$  in all three configurations) and so is essentially  $\varepsilon$ -invariant across all three tiers ( $\Delta \leq 0.003$ ); the attack’s minimum-norm objective dominates the budget constraint. The bottom line: increasing the adversarial budget either improves PRISM TPR (APGD-CE, APGD-DLR, Square) or leaves it statistically unchanged (FAB-T, where the residual  $\leq 0.3$  pp drifts sit well inside the  $\pm 1.4$  pp seed-pool 95% CI). The threat model conclusion is that the adversary gains nothing from spending more  $L_\infty$  budget against the detector.

## Q Recent Detector Baselines (2021): SID and SpectralDefense

The four matched-FPR baselines in §4.2 (LID, Mahalanobis, ODIN, Energy) predate 2021. To check that PRISM’s advantage holds against more recent detectors, we add two 2021 methods on the same CIFAR-10/ResNet-18 EVAL split, calibration scheme, and  $\varepsilon=8/255$  attack configs (5 seeds  $\times n=1000$ ):

- **SID** Tian et al. [2021] — sensitivity-inconsistency detector. We use the *unsupervised, no-training* variant: the score is the Jensen–Shannon divergence between  $\text{softmax}(f(x))$  and  $\text{softmax}(f(T(x)))$ , where  $T$  is a DCT low-pass reconstruction (keep fraction 0.70, selected on a held-out PGD probe). Threshold is fit on the clean split only, exactly as for LID/Mahalanobis/ODIN/Energy.
- **SpectralDefense** Harder et al. [2021] — a logistic-regression detector on the log-magnitude 2D-FFT, in both published variants: *InputMFS* (Fourier spectrum of the input image) and the stronger *LayerMFS* (Fourier spectrum of feature maps, channel-averaged per layer over layer2/3/4). Unlike PRISM and the other baselines, both are *supervised, attack-specific*: one classifier is trained per attack on (clean, adversarial) pairs from the disjoint REF split. They therefore see strictly more information than PRISM, which never trains on adversarials.

| Detector                                        | Type                 | FGSM         | PGD          | Square       | CW           |
|-------------------------------------------------|----------------------|--------------|--------------|--------------|--------------|
| SID Tian et al. [2021]                          | unsup. (no train)    | 0.296        | 0.519        | 0.286        | 0.601        |
| SpectralDefense Harder et al. [2021] (InputMFS) | superv. (per attack) | <b>1.000</b> | <b>0.999</b> | 0.208        | 0.244        |
| ↪ LayerMFS variant                              | superv. (per attack) | 0.861        | 0.899        | 0.711        | 0.458        |
| <b>PRISM (Full)</b>                             | unsup. (no train)    | 0.882        | 0.987        | <b>0.886</b> | <b>0.938</b> |

Table 29: Recent-detector comparison on CIFAR-10/ResNet-18 (5 seeds  $\times n=1000$ , design FPR  $\leq 0.10$ ). SpectralDefense numbers are the *in-distribution* diagonal (detector trained on the same attack it is evaluated on). **SID** (unsupervised) is weak across all four families (mean TPR 0.43), far below PRISM (0.92). **SpectralDefense/InputMFS** is near-perfect on gradient- $L_\infty$  (FGSM/PGD) but—even *trained on the target attack*—barely beats chance on Square (0.21) and CW (0.24): its input-Fourier signal captures dense gradient- $L_\infty$  perturbations but not localized black-box patches or minimal-norm  $L_2$  perturbations. The stronger **LayerMFS** variant (feature-map Fourier) substantially improves Square (0.21 $\rightarrow$ 0.71) and CW (0.24 $\rightarrow$ 0.46), confirming the gap is not an artifact of the weaker input-space variant—yet it still trails PRISM on both (0.71 vs 0.89; 0.46 vs 0.94), and is still supervised. PRISM is the only detector that holds  $\geq 0.88$  on every family, without supervision. (Bold marks the per-column best; the FGSM/PGD wins are supervised and in-distribution, not apples-to-apples with unsupervised PRISM.)

**Cross-attack transfer.** Because SpectralDefense trains one detector per attack, we can measure how well a detector trained on attack  $A$  transfers to attack  $B$  (Table 30, InputMFS). The two gradient- $L_\infty$  attacks (FGSM, PGD) transfer to each other almost perfectly ( $\geq 0.999$ ), confirming they share a Fourier signature. But a detector trained on FGSM or PGD collapses to a mean TPR of 0.17 when evaluated on Square or CW, and *no* training attack yields a detector that covers all four families. The stronger LayerMFS variant does *not* resolve this: its mean in-distribution (diagonal) TPR is higher (0.73 vs InputMFS 0.61), but its cross-attack (off-diagonal) mean is *lower* (0.38 vs 0.54)—i.e. LayerMFS is *more* attack-specific, collapsing 0.73 $\rightarrow$ 0.38 off its training attack. This is the supervised counterpart of PRISM’s design goal: a single, attack-agnostic detector that does not need to anticipate the attack.

**Protocol notes.** SID and SpectralDefense use the same three disjoint test-index ranges as the other baselines (REF 5000–5999, threshold-fit 6000–6999, EVAL 8000–9999). For tractability CW here uses a standard detection-benchmark configuration (bss=5, max\_iter=10, attack success rate 0.93 on the EVAL split) rather than the heavier canonical CW of Table 2; this does not affect the qualitative conclusion, since CW is poorly detected by both SpectralDefense variants and PRISM beats them regardless. We report both the *InputMFS* and the feature-space *LayerMFS* variants (LayerMFS features are the per-layer channel-averaged 2D-FFT log-magnitude of layer2/3/4 activations); LayerMFS data are in experiments/evaluation/recent\_baselines\_layermfs/.

| Trained on | FGSM         | Evaluated on |              |              |
|------------|--------------|--------------|--------------|--------------|
|            |              | PGD          | Square       | CW           |
| FGSM       | <u>1.000</u> | 0.999        | 0.139        | 0.182        |
| PGD        | 1.000        | <u>0.999</u> | 0.150        | 0.219        |
| Square     | 0.952        | 0.811        | <u>0.206</u> | 0.154        |
| CW         | 0.894        | 0.809        | 0.143        | <u>0.246</u> |

Table 30: SpectralDefense (InputMFS) cross-attack transfer, 5-seed mean L1 TPR. Underlined entries are in-distribution (train attack = eval attack). Gradient- $L_\infty$  detectors (FGSM/PGD rows) transfer to each other ( $\geq 0.999$ ) but collapse on Square/CW ( $\leq 0.22$ ); no row covers all four columns. Source: experiments/evaluation/recent\_baselines/ (per-seed JSONs, analyze\_cross\_attack.py).

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The four contributions in the abstract (TAMM, CADG, SACD, TAMSH) map one-to-one to Sections 3.1–3.4 (method) and 4.2–4.8 (experiments). All headline numbers (0.918/0.938/1.000 mean TPR, +33.6 pp baseline beat, 2.6-query campaign-detection latency, +26 pp recovery gap) are reproduced verbatim from the locked artifact JSONs listed in the experiments section.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Limitations are listed in Section 5 (computational cost of persistent homology, 150-point subsampling caveat, 39.78 ms (CIFAR-10) / 50.35 ms (CIFAR-100) mean GPU latency on an RTX 4090, conformal calibration assumes matched clean distribution, no guarantee against adaptive adversaries who target the topological score). Honest negative results are reported for the pure-topology TAMSH router (+10.6 pp vs. the +15 pp gate; see Appendix M) and the structurally-bounded SACD ASR-gap metric (Appendix K).

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.

- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The single theoretical result is Proposition 1 (split-conformal coverage). Its assumption (i.i.d. clean calibration scores) is stated in Section 3.2; a complete proof via the exchangeability argument of Vovk et al. [2005] appears in Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 4.1 specifies the backbone (CIFAR-native ResNet-18 trained from scratch via `scripts/pretrain_cifar_backbone.py`), per-layer weights, attack hyperparameters (FGSM  $\epsilon=8/255$ ; PGD-40 with 10 restarts; CW `max_iter=100`, `bss=9`; Square 5000 queries; AutoAttack standard ensemble), 5 fixed seeds (42, 123, 456, 789, 999),  $n=1000$  test images per seed, conformal  $\alpha \in \{0.10, 0.03, 0.005\}$ , and the exact Vast.ai launch scripts `run_vastai_full.sh` and `run_vastai_cifar100.sh`. Each result JSON is SHA-256 manifested for tamper-evident reproducibility.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.

- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: CIFAR-10 and CIFAR-100 are public benchmarks. An anonymized code release with the full Vast.ai launch scripts (`run_vastai_full.sh`, `run_vastai_cifar100.sh`), the resumable orchestrator, the SHA-256 artifact manifest, and all per-seed result JSONs accompanies this submission. The reproduction command is one line: `bash run_vastai_full.sh` on a CUDA 12 GPU instance with PyTorch 2.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).



- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4.1 specifies dataset splits (profile [0–4999], calibration [5000–6999], validation [7000–7999], evaluation [8000–9999] on the CIFAR-10 test set), conformal  $\alpha$  levels, attack hyperparameters, layer weights, backbone training schedule (SGD lr 0.1 with cosine annealing, 200 epochs, batch 256), and the seed list. Hyperparameter sensitivity over  $n_{\text{cal}}, n_{\text{sub}}, K, \lambda, p_{\text{alert}}, \delta, k_{\text{consec}}$  is reported in Appendix D.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All headline numbers are reported as mean  $\pm \sigma$  across 5 fixed seeds and accompanied by Wilson 95% binomial confidence intervals computed from pooled counts ( $n_{\text{adv}}, n_{\text{clean}}$  totals); see Tables 2–9. Per-channel ablation deltas are accompanied by paired  $t$ -tests (per-seed, per-attack), with  $p$ -values reported in the ablation aggregate JSON.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The full CIFAR-10 pipeline (backbone training, profile build, ensemble training, calibration, 5-seed evaluation across 5 attack families, ablation, campaign, recovery, baselines) runs in approximately one GPU-day on a single 24 GB-class GPU (Vast.ai 5090-class node); CIFAR-100 follows the same envelope. Post-fix local supplementary runs (baselines aggregation, CW re-runs at weakened parameters, uniform-router ablation) total ~33 minutes on a single GTX 1060.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The work is defensive adversarial-ML research using only public CIFAR-10/CIFAR-100 benchmarks; it involves no human subjects, no personally identifiable information, and no scraping. Reproductions of baseline attacks (FGSM/PGD/CW/Square/AutoAttack) and baseline detectors (LID/Mahalanobis/ODIN/Energy) use published, peer-reviewed implementations.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: A "Broader impact" paragraph in Section 5 states that PRISM is a defensive runtime monitor designed to improve the reliability of deployed ML systems under adversarial conditions; the principal anticipated impact is positive (safer deployment). We do not foresee misuse beyond standard adversarial-ML security research, since the system detects rather than generates adversarial inputs.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out

that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The released artifacts are a CIFAR-10/CIFAR-100 adversarial-detection model and accompanying calibration thresholds; they carry no significant dual-use risk distinct from the underlying public datasets or the standard ResNet-18 architecture. The paper does not release generative models, language models, or scraped data.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: CIFAR-10/CIFAR-100 Krizhevsky [2009] are cited with their canonical reference; both are MIT-style public benchmarks. Baseline detectors (LID Ma et al. [2018], Mahalanobis Lee et al. [2018], ODIN Liang et al. [2018], Energy Liu et al. [2020]) and attack implementations (FGSM, PGD, CW, Square, AutoAttack Croce and Hein [2020], Adversarial Robustness Toolbox Nicolae et al. [2018]) are cited with their original papers. Persistent-homology computation uses gudhi v3.12 GUDHI Project [2015] (MIT license). All licenses are respected.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.

- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

### 13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The anonymized supplemental code package contains the full PRISM source tree (`src/tamm`, `src/cadg`, `src/sacd`, `src/tamsh`), per-module docstrings, reproduction scripts, configuration files (`configs/default.yaml`, `configs/cifar100.yaml`), and a SHA-256 manifest of every result JSON. Trained backbone/calibrator/scorer/expert checkpoints are released under the same license as the code.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The work uses public image-classification benchmarks (CIFAR-10/CIFAR-100). No crowdsourcing, no human subjects, no annotation labor was involved.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: No human subjects research is conducted; the work is entirely on public image-classification benchmarks. IRB review is therefore not applicable.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not used in any core methodological component (TAMM persistence-diagram extraction, CADG conformal calibration, SACD CUSUM/BOCPD monitor, or TAMSH expert routing). Any writing or editing assistance falls under the policy exclusion for non-methodological usage.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.